

Mid-term Exam (Graph Neural Networks - Fall 2026)

해설지 ENG

Full Name:

Student ID:

Note: Students should write in English for evaluation.

1회차

1. (Relational Classification) Consider an undirected graph G of five nodes given in the following figure. Nodes 1 and 5 are labelled ($P(Y_1) = 1, P(Y_5) = 0$) and the others are unlabelled and initialized to 0.5. The relational classifier updates every unlabelled node as:

$$P(Y_v = c) = (\sum_{(v,u) \in E} A_{v,u} P(Y_u = c)) / (\sum_{(v,u) \in E} A_{v,u})$$




Figure. Graph G

- Explain the key assumption behind relational classification, and state one limitation.
- Calculate $P(Y_v)$ for every unlabelled node after the 1st iteration.
- Calculate $P(Y_v)$ after the 2nd iteration, using the previous iteration's values only.

Solution

ENG)

- Write these three sentences.
 - Relational classification assumes homophily: connected nodes tend to share the same label.
 - It repeatedly sets each unlabelled node to the weighted average of its neighbours' class probabilities.
 - Its limitation is that it uses only the graph structure and ignores node features.

Write the initial values: node 1 is 1, node 5 is 0, the rest start at 0.5.

First iteration. Only the unlabelled nodes are replaced by the mean of their neighbours.

Node 1 is labelled, so it never changes. $P = 1$

Node 2: neighbours 1, 3, 4. $(1 + 0.5 + 0.5)/3 = 0.6667$

Node 3: neighbours 2, 5. $(0.5 + 0)/2 = 0.25$

Node 4: neighbours 1, 2, 5. $(1 + 0.5 + 0)/3 = 0.5$

Node 5 is labelled, so it never changes. $P = 0$

b) $P(Y_1) = 1, P(Y_2) = 0.6667, P(Y_3) = 0.25, P(Y_4) = 0.5, P(Y_5) = 0$

Second iteration. Use only the values from the first iteration, not the ones just computed.

Node 1 is labelled, so it never changes. $P = 1$

Node 2: neighbours 1, 3, 4. $(1 + 0.25 + 0.5)/3 = 0.5833$

Node 3: neighbours 2, 5. $(0.6667 + 0)/2 = 0.3333$

Node 4: neighbours 1, 2, 5. $(1 + 0.6667 + 0)/3 = 0.5556$

Node 5 is labelled, so it never changes. $P = 0$

c) $P(Y1) = 1$, $P(Y2) = 0.5833$, $P(Y3) = 0.3333$, $P(Y4) = 0.5556$, $P(Y5) = 0$

2. (FastGCN) Consider the undirected graph G with adjacency matrix A and node feature vector $h^{(0)}$ below. FastGCN replaces the full neighbourhood sum of a GCN layer by a layer-wise Monte-Carlo estimate:

$$A = \begin{bmatrix} 0 & 1 & 1 & 1 & 1 \\ 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 1 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 \end{bmatrix}$$

$$h^{(0)} = \begin{bmatrix} 1 \\ 0 \\ 0 \\ -2 \\ 1 \end{bmatrix}$$

$$h^{(l+1)}(v) = \sigma \left(\frac{1}{t_l} \sum_{j=1}^{t_l} \tilde{A}(v, u_j) h^{(l)}(u_j) \right)$$

where $\tilde{A} = D^{-1}A$ and the t_l nodes are drawn once per layer from $q(u) \propto \|\tilde{A}(:, u)\|^2$, with $\text{ReLU}(x) = \max(0, x)$.

- How does layer-wise sampling differ from the node-wise sampling of GraphSAGE?
- Calculate the importance distribution $q(u)$ over the five nodes.
- Suppose the sampled set is $\{C, D\}$ ($t_l = 2$). Calculate the estimate of $h^{(1)}(A)$.

Solution

ENG)

a) Write these three sentences.

- In node-wise sampling, each node samples its own neighbours, so the cost grows exponentially with depth.
- FastGCN samples a fixed-size set of nodes once per layer, shared by all nodes in that layer.
- Therefore the cost grows only linearly with depth instead of exponentially.

Degrees: $d(A)=4$, $d(B)=2$, $d(C)=2$, $d(D)=3$, $d(E)=1$

Divide each row by its degree. $\tilde{A} = [0, 0.25, 0.25, 0.25, 0.25]; [0.5, 0, 0, 0.5, 0]; [0.5, 0, 0, 0.5, 0]; [0.3333, 0.3333, 0.3333, 0, 0]; [1, 0, 0, 0, 0]$ (each row sums to 1)

$q(u)$ uses the squared norm of column u of $\tilde{A}$, not of a row.

Column A: squared sum = 1.6111

Column B: squared sum = 0.1736

Column C: squared sum = 0.1736

Column D: squared sum = 0.5625

Column E: squared sum = 0.0625

Divide by the total 2.5833 to turn them into probabilities.

b) $q(A) = 0.6237$, $q(B) = 0.0672$, $q(C) = 0.0672$, $q(D) = 0.2177$, $q(E) = 0.0242$

For (c) look only at the sampled nodes: read $\tilde{A}(A, u)$ for each.

$\tilde{A}(A, C) h(C) = 0.25 \times 0 = 0$

$\tilde{A}(A, D) h(D) = 0.25 \times -2 = -0.5$

Add them and divide by the number sampled (2), not by the degree: $(0 + -0.5)/2 = -0.25$

ReLU clips the negatives: $-0.25 \rightarrow 0$

c) $h_1(A) = 0$

3. (ASGCN) Consider the undirected graph G with adjacency matrix A and node feature vector $h^{(0)}$ below. ASGCN rewrites a GCN layer as an expectation:

$$A = \begin{bmatrix} 0 & 1 & 0 & 1 & 0 \\ 1 & 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 0 & 1 \\ 1 & 1 & 0 & 0 & 1 \\ 0 & 1 & 1 & 1 & 0 \end{bmatrix}$$

$$h^{(0)} = \begin{bmatrix} -2 \\ -1 \\ 3 \\ 2 \\ 4 \end{bmatrix}$$

$$h^{(l+1)}(v_i) = \sigma(W^{(l)} N(v_i) E_{p(u_j | v_i)}[h^{(l)}(u_j)])$$

where $\tilde{A} = D^{-1}A$, $N(v_i) = \sum_j \tilde{A}_{ij}$, $p(u_j | v_i) = \tilde{A}_{ij} / N(v_i)$, $W^{(l)} = 1$ and $\text{ReLU}(x) = \max(0, x)$.

- What does the proposal distribution depend on in ASGCN that it does not in FastGCN?
- Calculate $N(v_i)$ and $p(u_j | v_i)$ for the parent node $v_i = E$.
- Calculate $E_{p(u_j | E)}[h^{(0)}(u_j)]$ and the resulting $h^{(1)}(E)$.

Solution

ENG)

- Write these three sentences.

- In FastGCN the sampling distribution is fixed in advance and does not depend on the parent nodes.
- In ASGCN the proposal distribution is conditioned on the nodes of the upper layer.
- This makes almost every sampled node useful and reduces the variance of the estimate.

Degrees: $d(A)=2$, $d(B)=3$, $d(C)=1$, $d(D)=3$, $d(E)=3$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0, 0.5, 0]; [0.3333, 0, 0, 0.3333, 0.3333]; [0, 0, 0, 0, 1]; [0.3333, 0.3333, 0, 0, 0.3333]; [0, 0.3333, 0.3333, 0.3333, 0]$ (each row sums to 1)

The parent node is E, so only its row of $\tilde{A}$ matters.

$N(E) = \text{sum of that row} = 1$. With row normalisation it is 1.

$$p(B|E) = \tilde{A}(E,B) / N = 0.3333$$

$$p(C|E) = \tilde{A}(E,C) / N = 0.3333$$

$$p(D|E) = \tilde{A}(E,D) / N = 0.3333$$

$$b) N(E) = 1, p(B|E) = 0.3333, p(C|E) = 0.3333, p(D|E) = 0.3333$$

The expectation is the sum of probability times value: $0.3333 \times -1 + 0.3333 \times 3 + 0.3333 \times 2 = 1.3333$

$N = 1$ and $W = 1$, so the value is unchanged. It equals the plain neighbour mean.

No negative value, so ReLU changes nothing.

$$c) E = 1.3333, h_1(E) = 1.3333$$

4. (APPNP) Given a graph with an adjacency matrix A and a prediction matrix H produced by an MLP as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{bmatrix}$$

$$H = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$$

APPNP propagates the predictions with a teleport probability α :

$$Z^{(0)} = H, \quad Z^{(k+1)} = (1 - \alpha) \tilde{A} Z^{(k)} + \alpha H$$

where $\tilde{A} = D^{-1}A$ and $\alpha = 0.2$. No nonlinearity is applied between steps.

- Explain the teleport mechanism and why it prevents over-smoothing.
- Calculate $\tilde{A}$ and $Z^{(1)}$.
- Calculate $Z^{(2)}$.

Solution

ENG)

a) Write these three sentences.

- At every propagation step the walk returns to the root node with probability alpha. This is the teleport mechanism.
- Because a fixed fraction of the mass is pulled back at every step, the distribution stays local.
- So we can use many propagation steps without over-smoothing, and propagation has no extra parameters.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=2$, $d(D)=2$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0.5, 0]; [0.5, 0, 0, 0.5]; [0.5, 0, 0, 0.5]; [0, 0.5, 0.5, 0]$ (each row sums to 1)

b) $\tilde{A} = [0, 0.5, 0.5, 0]; [0.5, 0, 0, 0.5]; [0.5, 0, 0, 0.5]; [0, 0.5, 0.5, 0]$

Start from $Z_0 = H = [0, 1, 0, 0]$

Propagate one step: $\tilde{A} Z_0 = [0.5, 0, 0, 0.5]$

Multiply by $(1-\alpha) = 0.8$: $[0.4, 0, 0, 0.4]$

Add $\alpha H = [0, 0.2, 0, 0]$. This is the teleport term.

$$Z_1 = [0.4, 0.2, 0, 0.4]$$

One more step: $\tilde{A} Z_1 = [0.1, 0.4, 0.4, 0.1]$

Multiply by $(1-\alpha)$ and add αH **again**. It is added at every step.

c) $Z_2 = [0.08, 0.52, 0.32, 0.08]$

5. (GPR-GNN) Given a graph with an adjacency matrix A and hidden features $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 4 \\ -2 \\ 2 \\ 2 \end{bmatrix}$$

GPR-GNN separates feature extraction from topological propagation:

$$H^{(k)} = \tilde{A} H^{(k-1)}, \quad Z = \sum_{k=0}^K \gamma_k H^{(k)}$$

where $\tilde{A} = D^{-1}A$, $K = 2$, $\gamma_0 = 0.7$, $\gamma_1 = 0.3$, $\gamma_2 = -0.2$.

- What can GPR-GNN do that APPNP cannot, and why do negative γ_k matter?
- Calculate $H^{(1)}$ and $H^{(2)}$.
- Calculate Z .

Solution

ENG)

- Write these three sentences.
 - APPNP fixes the coefficients to alpha times (1-alpha) to the power k, which are always positive and decreasing.
 - GPR-GNN learns the weights, and they are allowed to be negative.
 - A negative weight subtracts that hop, so the model can act as a high-pass filter on heterophilic graphs.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=2$, $d(D)=2$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0.5, 0]; [0.5, 0, 0, 0.5]; [0.5, 0, 0, 0.5]; [0, 0.5, 0.5, 0]$ (each row sums to 1)

H_0 stays as it is: $[4, -2, 2, 2]$

$H_1 = \tilde{A} H_0 = [0, 3, 3, 0]$

$H_2 = \tilde{A} H_1 = [3, 0, 0, 3]$

b) $H_1 = [0, 3, 3, 0]$, $H_2 = [3, 0, 0, 3]$

Now take the weighted sum with gamma. There is no ReLU during propagation.

$\gamma_0 H_0 = 0.7 \times [4, -2, 2, 2] = [2.8, -1.4, 1.4, 1.4]$

$\gamma_1 H_1 = 0.3 \times [0, 3, 3, 0] = [0, 0.9, 0.9, 0]$

$\gamma_2 H_2 = -0.2 \times [3, 0, 0, 3] = [-0.6, 0, 0, -0.6]$

γ_2 is negative, so that hop is subtracted. Dropping the sign gives a wrong answer.

Add the three vectors element-wise.

c) $Z = [2.2, -0.5, 2.3, 0.8]$

6. (MixHop) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 1 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 1 \\ -3 \\ 3 \\ -1 \end{bmatrix}$$

A MixHop layer concatenates several powers of the adjacency matrix:

$$H^{(i+1)} = \parallel_{j \in P} \sigma(\hat{A}^j H^{(i)} W_j^{(i)})$$

where $P = \{0, 1, 2\}$, $\hat{A} = D^{-1}A$, $W_j^{(i)} = 1$ and $\text{ReLU}(x) = \max(0, x)$.

- State Theorem 1 of MixHop about the two-hop Delta Operator.
- Calculate $\hat{A}^0 H^{(0)}$, $\hat{A}^1 H^{(0)}$ and $\hat{A}^2 H^{(0)}$.
- Calculate $H^{(1)}$ and the Delta Operator $\sigma(\hat{A} H^{(0)}) - \sigma(\hat{A}^2 H^{(0)})$.

Solution

ENG)

a) Write these three sentences.

- Theorem 1 says that a vanilla GCN cannot represent the two-hop Delta Operator, even with infinitely many layers.
- Stacking GCN layers collapses to $\hat{A}$ to the power l times W , so the intermediate hops are never kept separately.
- MixHop computes each power separately and concatenates them, so a later layer can subtract one from another.

Degrees: $d(A)=1$, $d(B)=3$, $d(C)=2$, $d(D)=2$

Divide each row by its degree. $\tilde{A} = [0, 1, 0, 0]; [0.3333, 0, 0.3333, 0.3333]; [0, 0.5, 0, 0.5]; [0, 0.5, 0.5, 0]$ (each row sums to 1)

$\hat{A}^0$ is the identity, so it is just $H^0: [1, -3, 3, -1]$

$\hat{A}^1 H^0 = [-3, 1, -2, 0]$ (one hop, the neighbour mean)

$\hat{A}^2 H^0 = \tilde{A}(\tilde{A} H^0) = [1, -1.6667, 0.5, -0.5]$ (two hops)

ReLU is applied to each block separately, not after concatenation.

After ReLU: $[1, 0, 3, 0] \parallel [0, 1, 0, 0] \parallel [1, 0, 0.5, 0]$

The delta operator subtracts the two-hop block from the one-hop block.

b) $A^0 H = [1, -3, 3, -1]$, $A^1 H = [-3, 1, -2, 0]$, $A^2 H = [1, -1.6667, 0.5, -0.5]$

c) $H^1 =$ concatenation of the three blocks: $[1, 0, 3, 0] \parallel [0, 1, 0, 0] \parallel [1, 0, 0.5, 0]$

Delta = $[-1, 1, -0.5, 0]$

7. (**PairNorm**) After one GCN layer the node features of a four-node graph are

$$X = \begin{bmatrix} 2 & -2 \\ 4 & -1 \\ 4 & 3 \\ 3 & 2 \end{bmatrix}$$

PairNorm normalizes them with a centering step and a scaling step:

$$x_i^c = x_i - (1/n) \sum_i x_i \quad (\text{Center})$$

$$x_i = s \cdot x_i^c / \text{sqrt}((1/n) \sum_i \|x_i^c\|_2^2) \quad (\text{Scale})$$

with $s = 1$ and $n = 4$.

- Explain node-wise and feature-wise over-smoothing, and what PairNorm keeps constant.
- Calculate the centered features.
- Calculate the scaling denominator and the final x_i .

Solution

ENG)

- Write these three sentences.
 - Node-wise over-smoothing means the features of connected nodes become too similar.
 - Feature-wise over-smoothing means all node features become indistinguishable across the whole graph.
 - PairNorm keeps the total pairwise squared distance between node representations constant across layers.

First take the mean over the nodes, separately for each feature column.

Column 1: $(2 + 4 + 4 + 3)/4 = 3.25$

Column 2: $(-2 + -1 + 3 + 2)/4 = 0.5$

Subtract the mean from every row. Check: each column now sums to zero.

Node 1: squared norm = 7.8125

Node 2: squared norm = 2.8125

Node 3: squared norm = 6.8125

Node 4: squared norm = 2.3125

Mean = $(7.8125 + 2.8125 + 6.8125 + 2.3125)/4 = 4.9375$

Take the square root: $\text{sqrt}(4.9375) = 2.222$. Divide every row by this single number.

b) Mean = $[3.25, 0.5]$, centered = $[-1.25, -2.5]; [0.75, -1.5]; [0.75, 2.5]; [-0.25, 1.5]$

c) Scaling denominator = 2.222, $x = [-0.5625, -1.1251]; [0.3375, -0.6751]; [0.3375, 1.1251]; [-0.1125, 0.6751]$

8. (DropEdge) Consider the complete graph K_4 on nodes A, B, C and D with the adjacency matrix A and initial node feature matrix $H^{(0)}$ below.

$$A = \begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 3 \\ -3 \\ 3 \\ 2 \end{bmatrix}$$

DropEdge keeps each undirected edge independently with probability $1 - p$ at every epoch:

$$H^{(1)} = \sigma((D^{(t)})^{-1} A^{(t)} H^{(0)})$$

where $A^{(t)}$ and $D^{(t)}$ are rebuilt from the surviving edges only, $p = 1/3$, and $\text{ReLU}(x) = \max(0, x)$.

- Which two problems does DropEdge target, and how does it differ from Dropout?
- Calculate the expected number of surviving edges, and $H^{(1)}$ on the full graph.
- Suppose the edges (B, C) and (A, C) are dropped. Recalculate $H^{(1)}$.

Solution

ENG)

- Write these three sentences.
 - DropEdge targets over-fitting and over-smoothing.
 - A different random subgraph is used at every epoch, which works as data augmentation on the graph.
 - Dropout zeroes entries of the feature tensor, while DropEdge removes edges of the adjacency matrix.

K_4 has 6 edges. Each survives with probability $1 - 1/3 = 2/3$, so the expectation is 4.

First the full graph: every node has degree 3.

A: neighbours B, C, D. $(-3 + 3 + 2)/3 = 0.6667$

B: neighbours A, C, D. $(3 + 3 + 2)/3 = 2.6667$

C: neighbours A, B, D. $(3 + -3 + 2)/3 = 0.6667$

D: neighbours A, B, C. $(3 + -3 + 3)/3 = 1$

Now drop the two given edges and **recount the degrees**. This is where most mistakes happen.

Degrees after dropping: $d(A)=2$, $d(B)=2$, $d(C)=1$, $d(D)=3$

A: neighbours B, D. $(-3 + 2)/2 = -0.5$

B: neighbours A, D. $(3 + 2)/2 = 2.5$

C: neighbours D. $(2)/1 = 2$

D: neighbours A, B, C. $(3 + -3 + 3)/3 = 1$

b) Expected number of surviving edges = 4, $H_1 = [0.6667, 2.6667, 0.6667, 1]$

c) $H_1 = [0, 2.5, 2, 1]$

9. (Node2vec) Consider an undirected graph G given in the following figure. A random walk has just moved from node A to node B , with $p = 2$ and $q = 0.5$.




Figure 1. Graph G

The unnormalized transition probability from B to a next node x is $\pi_{Bx} = \alpha_{pq}(A, x)$, where $\alpha = 1/p$ if $d(A, x) = 0$, $\alpha = 1$ if $d(A, x) = 1$, and $\alpha = 1/q$ if $d(A, x) = 2$.

- Explain BFS and DFS, and explain which one a small q makes the walk behave like.
- List all possible next nodes from node B and calculate π_{Bx} .
- Calculate the normalized transition probabilities.

Solution

ENG)

- Write these three sentences.
 - BFS explores the immediate neighbourhood first, so the embedding captures the structural role of a node.
 - DFS moves far away from the source, so the embedding captures community membership.
 - A small q makes $1/q$ large, so the walk moves outward and behaves like DFS.

The previous node is A and the walk is now at B . All distances are measured from A .

A : $x = A$ is the previous node, $d = 0$, $\alpha = 1/p = 0.5$

C : x is two hops from A , $d = 2$, $\alpha = 1/q = 2$

E : x is two hops from A , $d = 2$, $\alpha = 1/q = 2$

Sum them: $0.5 + 2 + 2 = 4.5$

b) $\pi(B \rightarrow A) = 0.5$, $\pi(B \rightarrow C) = 2$, $\pi(B \rightarrow E) = 2$

c) $P(A) = 0.1111$, $P(C) = 0.4444$, $P(E) = 0.4444$

10. (LINE) Consider an undirected graph G with node embeddings

$$u_A = 1, u_B = 0.5, u_C = -0.5, u_D = -1$$

and context embeddings

$$u'_A = 1, u'_B = 0.5, u'_C = -1, u'_D = -0.5.$$

The second-order proximity is $p_2(j|x) = \exp(u'_j u_x) / \sum_k \exp(u'_k u_x)$. Given $e^{0.25} = 1.28$, $e^{0.5} = 1.65$, $e^{-0.25} = 0.78$, $e^{-0.5} = 0.61$.

- What is second-order proximity, and how does it differ from first-order proximity?
- Calculate $p_2(A|C)$, $p_2(B|C)$, $p_2(C|C)$ and $p_2(D|C)$.

Solution

ENG)

- Write these three sentences.

- First-order proximity is the similarity between two nodes that are directly connected by an edge.
- Second-order proximity is the similarity between two nodes that share many neighbours, even without an edge.
- So LINE gives each node two vectors: its own embedding and its context embedding.

The base node is C with $u(C) = -0.5$. The other side uses the context vector u' .

$$A: u'(A) u(C) = 1 \times -0.5 = -0.5, \exp = 0.61$$

$$B: u'(B) u(C) = 0.5 \times -0.5 = -0.25, \exp = 0.78$$

$$C: u'(C) u(C) = -1 \times -0.5 = 0.5, \exp = 1.65$$

$$D: u'(D) u(C) = -0.5 \times -0.5 = 0.25, \exp = 1.28$$

The denominator is the sum of all four: $0.61 + 0.78 + 1.65 + 1.28 = 4.32$

$$b) p_2(A|C) = 0.1412, p_2(B|C) = 0.1806, p_2(C|C) = 0.3819, p_2(D|C) = 0.2963$$

11. (GCN) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 \\ 1 & 1 & 0 & 1 & 1 \\ 1 & 1 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 2 \\ 0 \\ 2 \\ 3 \\ -2 \end{bmatrix}$$

Assume that the hidden layer of a GCN model of all nodes at layer (k) can be calculated as:

$$H^{(k)} = \sigma(\tilde{A} \cdot H^{(k-1)}),$$

where $\tilde{A} = D^{-1}A$ and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- Explain how information is propagated in a GCN layer.
- Calculate $\tilde{A}$.
- Calculate the output of the GCN model at layer $k = 1$.

Solution

ENG)

a) Write these three sentences.

- In one layer each node aggregates the features of its direct neighbours and takes the average.
- The row-normalized matrix $\tilde{A}$ equals $D^{-1}A$, which performs this averaging.
- After k layers a node has received information from all nodes within k hops.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=4$, $d(D)=3$, $d(E)=1$

Divide each row by its degree. $\tilde{A} = [0, 0, 0.5, 0.5, 0]; [0, 0, 0.5, 0.5, 0]; [0.25, 0.25, 0, 0.25, 0.25]; [0.3333, 0.3333, 0, 0, 0]; [0, 0, 1, 0, 0]$ (each row sums to 1)

b) $\tilde{A} = [0, 0, 0.5, 0.5, 0]; [0, 0, 0.5, 0.5, 0]; [0.25, 0.25, 0, 0.25, 0.25]; [0.3333, 0.3333, 0.3333, 0, 0]; [0, 0, 1, 0, 0]$

Multiply $\tilde{A}$ by H^0 . Each entry is the mean of that node's neighbours.

A: neighbours C, D. $(2 + 3)/2 = 2.5$

B: neighbours C, D. $(2 + 3)/2 = 2.5$

C: neighbours A, B, D, E. $(2 + 0 + 3 + -2)/4 = 0.75$

D: neighbours A, B, C. $(2 + 0 + 2)/3 = 1.3333$

E: neighbours C. $(2)/1 = 2$

No negative value, so ReLU changes nothing.

c) $H_1 = [2.5, 2.5, 0.75, 1.3333, 2]$

12. (GraphSAGE) Consider an undirected graph G of five nodes given in the following figure. Each node has initial features that are the numbers standing next to it. According to the GraphSAGE model with an AGGREGATE that is a MEAN function, the feature of a node i at layer k can be updated as:

$$h_{N(i)}^{(k)} = \text{MEAN}(\{ h_u^{(k-1)}, \forall u \in N(i) \})$$

$$h_i^{(k)} = \text{ReLU}(h_i^{(k-1)} \parallel h_{N(i)}^{(k)})$$

where $\parallel$ is a concatenation and $\text{ReLU}(x) = \max(0, x)$.




Figure. Graph G

- What is the main idea behind GraphSAGE, and how does it differ from traditional GCNs?
- Calculate the feature of each node at $k = 1$.
- Calculate a graph-level embedding h_G by using a 'Mean' global pooling when $k = 1$.

Solution

ENG)

a) Write these three sentences.

- GraphSAGE learns an aggregator function instead of a fixed embedding for each node.
- It samples a fixed number of neighbours, so it can be trained in mini-batches on very large graphs.
- A GCN is transductive, but GraphSAGE is inductive, so it can embed a node that was never seen in training.

First compute the neighbour mean. The node itself is not included.

A: neighbours C, D. $(-1 + 2)/2 = 0.5$

B: neighbours C, D. $(-1 + 2)/2 = 0.5$

C: neighbours A, B, D, E. $(3 + 4 + 2 + 4)/4 = 3.25$

D: neighbours A, B, C. $(3 + 4 + -1)/3 = 2$

E: neighbours C. $(-1)/1 = -1$

Concatenate the node's own value with the neighbour mean. This is not an addition.

ReLU clips the negatives: $-1 \rightarrow 0$, $-1 \rightarrow 0$

The graph embedding is the mean per position: first $12 / 5$, second $5.25 / 5$

b) $h(A) = [3, 0.5]$, $h(B) = [4, 0.5]$, $h(C) = [0, 3.25]$, $h(D) = [2, 2]$, $h(E) = [4, 0]$

c) $h_G = [2.6, 1.25]$

13. (ClusterGCN) Consider an undirected graph G of six nodes given in the following figure. The graph G contains two clusters G_1 and G_2 . Each node has initial features that are the numbers standing next to it. According to the ClusterGCN model, the feature of a node i at layer k can be updated as:

$$h_{N(i)}^{(k)} = \text{MEAN}(\{ h_u^{(k-1)}, \forall u \in N(i), G_u = G_i \})$$

$$h_i^{(k)} = \text{ReLU}(h_i^{(k-1)} \parallel h_{N(i)}^{(k)})$$

where $\parallel$ is a concatenation and $\text{ReLU}(x) = \max(0, x)$.




Figure. Graph G

- What is mini-batch training, and why does ClusterGCN need it?
- Calculate the output representations of all nodes at layer $k = 1$.

Solution

ENG)

- Write these three sentences.
 - Mini-batch training updates the parameters using a small subset of the data at a time.
 - A large graph does not fit in memory, so ClusterGCN partitions the graph into densely connected clusters.
 - A good partition keeps most edges inside a batch, so little information is lost at the boundary.

Edges that cross a group are treated as absent. Cross them out first.

A: neighbours C. $(-1)/1 = -1$

B: neighbours C. $(-1)/1 = -1$

C: neighbours A, B. $(2 + 0)/2 = 1$

D: neighbours E, F. $(3 + 1)/2 = 2$

E: neighbours D, F. $(4 + 1)/2 = 2.5$

F: neighbours D, E. $(4 + 3)/2 = 3.5$

Concatenate the neighbour mean after the node's own value, then apply ReLU to both entries.

b) $h(A) = [2, 0]$, $h(B) = [0, 0]$, $h(C) = [0, 1]$, $h(D) = [4, 2]$, $h(E) = [3, 2.5]$, $h(F) = [1, 3.5]$

14. (GraphSAINT) Consider an undirected graph G of six nodes given in the following figure. The graph G has a sampled subgraph G_1 . Each node has initial features that are the numbers standing next to it. According to the GraphSAINT model, the feature of a node i at layer k can be updated as:

$$h_{N(i)}^{(k)} = \text{MEAN}(\{ h_u^{(k-1)}, \forall u \in N(i), G_u = G_i \})$$

$$h_i^{(k)} = \text{ReLU}(h_i^{(k-1)} \parallel h_{N(i)}^{(k)})$$

where $\parallel$ is a concatenation and $\text{ReLU}(x) = \max(0, x)$.




Figure. Graph G

- How does GraphSAINT's sampling differ from node-wise or layer-wise sampling?
- Calculate the output representations of nodes A, B, C and D at layer $k = 1$.

Solution

ENG)

- Write these three sentences.
 - GraphSAGE and FastGCN sample while building the layers, so sampling is repeated at every layer.
 - GraphSAINT samples a subgraph first and then builds the full GNN on that subgraph.
 - So the cost does not grow with depth, and normalization coefficients correct the sampling bias.

Edges that leave the sampled subgraph are discarded.

A: neighbours C, D. $(2 + 3)/2 = 2.5$

B: neighbours D. $(3)/1 = 3$

C: neighbours A. $(-1)/1 = -1$

D: neighbours A, B. $(-1 + 4)/2 = 1.5$

Concatenate the neighbour mean after the node's own value, then apply ReLU.

b) $h(A) = [0, 2.5]$, $h(B) = [4, 3]$, $h(C) = [2, 0]$, $h(D) = [3, 1.5]$

15. (JK Network) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 1 & 0 \\ 1 & 0 & 0 & 1 & 1 \\ 1 & 0 & 0 & 1 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} -2 \\ 2 \\ -3 \\ 2 \\ 1 \end{bmatrix}$$

Assume that the output of a JK network model of all nodes at layer (k) can be calculated as:

$$H^{(k)} = \max(\sigma(\tilde{A} H^{(0)}), \sigma(\tilde{A} H^{(1)}), \dots, \sigma(\tilde{A} H^{(k-1)}))$$

where $\tilde{A} = D^{-1}A$ and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- Explain how Jumping Knowledge Networks help mitigate the over-smoothing problem.
- Calculate $\tilde{A}$ and $H^{(1)}$.
- Calculate the output representations at layer $k = 2$.

Solution

ENG)

a) Write these three sentences.

- In a deep GCN every node repeatedly averages its neighbours, so all representations become similar. This is over-smoothing.
- JKNet keeps the representation of every layer and combines them at the end by max, concatenation or an LSTM.
- The best depth is different for each node, so each node can select the depth that suits it.

Degrees: $d(A)=3$, $d(B)=3$, $d(C)=2$, $d(D)=3$, $d(E)=1$

Divide each row by its degree. $\text{Atilde} = [0, 0.3333, 0.3333, 0.3333, 0]; [0.3333, 0, 0, 0.3333, 0.3333]; [0.5, 0, 0, 0.5, 0]; [0.3333, 0.3333, 0.3333, 0, 0]; [0, 1, 0, 0, 0]$ (each row sums to 1)

b) $\text{Atilde} = [0, 0.3333, 0.3333, 0.3333, 0]; [0.3333, 0, 0, 0.3333, 0.3333]; [0.5, 0, 0, 0.5, 0]; [0.3333, 0.3333, 0.3333, 0, 0]; [0, 1, 0, 0, 0]$

First build $H_1 = \text{ReLU}(\text{Atilde } H_0)$

A: neighbours B, C, D. $(2 + -3 + 2)/3 = 0.3333$

B: neighbours A, D, E. $(-2 + 2 + 1)/3 = 0.3333$

C: neighbours A, D. $(-2 + 2)/2 = 0$

D: neighbours A, B, C. $(-2 + 2 + -3)/3 = -1$

E: neighbours B. $(2)/1 = 2$

ReLU clips the negatives: $-1 \rightarrow 0$

$H_1 = [0.3333, 0.3333, 0, 0, 2]$

For $k = 2$ there are two terms. First: $\text{ReLU}(\text{Atilde } H_0) = [0.3333, 0.3333, 0, 0, 2]$

Second: $\text{ReLU}(\text{Atilde } H_1) = [0.1111, 0.7778, 0.1667, 0.2222, 0.3333]$

The max is taken element-wise, not by choosing one whole vector.

c) $H_2 = [0.3333, 0.7778, 0.1667, 0.2222, 2]$

16. (GCNII) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 \\ 1 & 1 & 0 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 3 \\ -1 \\ 2 \\ -3 \\ 0 \end{bmatrix}$$

Assume that the output of a GCNII model of all nodes at layer (k) can be calculated as:

$$H^{(k)} = \sigma[((1 - \beta) I_n) \cdot ((1 - \alpha) \tilde{A} \cdot H^{(k-1)} + \alpha H^{(0)})]$$

where $\tilde{A} = D^{-1}A$, I_n is the identity matrix, $\alpha = 0.2$, $\beta = 0.5$, and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- What are the two major components introduced in GCNII to overcome the depth limitation?
- Calculate $\tilde{A}$.
- Calculate the output representations at layer $k = 1$.

Solution

ENG)

- Write these three sentences.
 - The two components are the initial residual connection and the identity mapping.
 - The initial residual connection adds a fraction alpha of the input H zero back at every layer.
 - The identity mapping keeps the weight matrix close to the identity, so a deep model is not destroyed.

Degrees: $d(A)=3$, $d(B)=1$, $d(C)=2$, $d(D)=3$, $d(E)=3$

Divide each row by its degree. $\tilde{A} = [0, 0, 0.3333, 0.3333, 0.3333]; [0, 0, 0, 0, 1]; [0.5, 0, 0, 0.5, 0]; [0.3333, 0, 0.3333, 0, 0.3333]; [0.3333, 0.3333, 0, 0.3333, 0]$ (each row sums to 1)

b) $\tilde{A} = [0, 0, 0.3333, 0.3333, 0.3333]; [0, 0, 0, 0, 1]; [0.5, 0, 0, 0.5, 0]; [0.3333, 0, 0.3333, 0, 0.3333]; [0.3333, 0.3333, 0, 0.3333, 0]$

Start with the inner bracket. $\tilde{A} H^0 = [-0.3333, 0, 0, 1.6667, -0.3333]$

Multiply by $(1-\alpha) = 0.8$: $[-0.2667, 0, 0, 1.3333, -0.2667]$

Add $\alpha H^0 = 0.2 \times [3, -1, 2, -3, 0] = [0.6, -0.2, 0.4, -0.6, 0]$

Inner bracket = $[0.3333, -0.2, 0.4, 0.7333, -0.2667]$

Multiply by the outer $(1-\beta) = 0.5$. Dropping this factor scales the answer.

ReLU clips the negatives: $-0.1 \rightarrow 0$, $-0.1333 \rightarrow 0$

c) $H^1 = [0.1667, 0, 0.2, 0.3667, 0]$

17. (DeepGCN) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 0 & 1 & 1 \\ 0 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} -1 \\ 3 \\ 3 \\ 1 \end{bmatrix}$$

Assume that the output of a DeepGCN model with a residual connection at layer (k) is:

$$H^{(k)} = \sigma(\tilde{A} \cdot H^{(k-1)}) + H^{(k-1)}$$

where $\tilde{A} = D^{-1}A$ and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- Explain why a residual connection allows a GCN to be made deeper.
- Calculate the output representations at layer $k = 1$ and $k = 2$.
- Compare $H^{(1)}$ with the output of a plain GCN without the residual connection.

Solution

ENG)

- Write these three sentences.
 - Deep GCNs suffer from vanishing gradients and over-smoothing, so adding layers makes the model worse.
 - A residual connection passes the input forward unchanged, so each layer only has to learn the change.
 - If no change is useful the layer can output zero, and gradients also flow back through the skip path.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=3$, $d(D)=3$

Divide each row by its degree. $\tilde{A} = [0, 0, 0.5, 0.5]; [0, 0, 0.5, 0.5]; [0.3333, 0.3333, 0, 0.3333]; [0.3333, 0.3333, 0.3333, 0]$ (each row sums to 1)

$\tilde{A} H^0 = [2, 2, 1, 1.6667]$

Apply ReLU first: $[2, 2, 1, 1.6667]$

Then add H^0 . The addition sits outside the ReLU, so negatives can survive.

For $k = 2$, feed H^1 back in and add H^1 , not H^0 .

$\tilde{A} H^1 = [3.3333, 3.3333, 2.8889, 3.3333]$, after ReLU $[3.3333, 3.3333, 2.8889, 3.3333]$

b) $H^1 = [1, 5, 4, 2.6667]$

$H^2 = [4.3333, 8.3333, 6.8889, 6]$

c) Plain GCN without residual: $H^1 = [2, 2, 1, 1.6667]$ (the node's own value and negatives disappear)

2회차

1. (Relational Classification) Consider an undirected graph G of five nodes given in the following figure. Nodes 1 and 5 are labelled ($P(Y_1) = 1$, $P(Y_5) = 0$) and the others are unlabelled and initialized to 0.5. The relational classifier updates every unlabelled node as:

$$P(Y_v = c) = (\sum_{(v,u) \in E} A_{v,u} P(Y_u = c)) / (\sum_{(v,u) \in E} A_{v,u})$$




Figure. Graph G

- Explain the key assumption behind relational classification, and state one limitation.
- Calculate $P(Y_v)$ for every unlabelled node after the 1st iteration.
- Calculate $P(Y_v)$ after the 2nd iteration, using the previous iteration's values only.

Solution

ENG)

- Write these three sentences.
 - Relational classification assumes homophily: connected nodes tend to share the same label.
 - It repeatedly sets each unlabelled node to the weighted average of its neighbours' class probabilities.
 - Its limitation is that it uses only the graph structure and ignores node features.

Write the initial values: node 1 is 1, node 5 is 0, the rest start at 0.5.

First iteration. Only the unlabelled nodes are replaced by the mean of their neighbours.

Node 1 is labelled, so it never changes. $P = 1$

Node 2: neighbours 1, 3, 4. $(1 + 0.5 + 0.5)/3 = 0.6667$

Node 3: neighbours 1, 2. $(1 + 0.5)/2 = 0.75$

Node 4: neighbours 1, 2. $(1 + 0.5)/2 = 0.75$

Node 5 is labelled, so it never changes. $P = 0$

b) $P(Y_1) = 1$, $P(Y_2) = 0.6667$, $P(Y_3) = 0.75$, $P(Y_4) = 0.75$, $P(Y_5) = 0$

Second iteration. Use only the values from the first iteration, not the ones just computed.

Node 1 is labelled, so it never changes. $P = 1$

Node 2: neighbours 1, 3, 4. $(1 + 0.75 + 0.75)/3 = 0.8333$

Node 3: neighbours 1, 2. $(1 + 0.6667)/2 = 0.8333$

Node 4: neighbours 1, 2. $(1 + 0.6667)/2 = 0.8333$

Node 5 is labelled, so it never changes. $P = 0$

c) $P(Y_1) = 1$, $P(Y_2) = 0.8333$, $P(Y_3) = 0.8333$, $P(Y_4) = 0.8333$, $P(Y_5) = 0$

2. (FastGCN) Consider the undirected graph G with adjacency matrix A and node feature vector $h^{(0)}$ below. FastGCN replaces the full neighbourhood sum of a GCN layer by a layer-wise Monte-Carlo estimate:

$$A = \begin{bmatrix} 0 & 1 & 1 & 1 & 0 \\ 1 & 0 & 0 & 1 & 1 \\ 1 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 1 \\ 0 & 1 & 0 & 1 & 0 \end{bmatrix}$$

$$h^{(0)} = \begin{bmatrix} -2 \\ -2 \\ 3 \\ 1 \\ -1 \end{bmatrix}$$

$$h^{(l+1)}(v) = \sigma \left(\frac{1}{t_l} \sum_{j=1}^{t_l} \tilde{A}(v, u_j) h^{(l)}(u_j) \right)$$

where $\tilde{A} = D^{-1}A$ and the t_l nodes are drawn once per layer from $q(u) \propto \|\tilde{A}(:, u)\|^2$, with $\text{ReLU}(x) = \max(0, x)$.

- How does layer-wise sampling differ from the node-wise sampling of GraphSAGE?
- Calculate the importance distribution $q(u)$ over the five nodes.
- Suppose the sampled set is $\{D, E\}$ ($t_l = 2$). Calculate the estimate of $h^{(1)}(A)$.

Solution

ENG)

a) Write these three sentences.

- In node-wise sampling, each node samples its own neighbours, so the cost grows exponentially with depth.
- FastGCN samples a fixed-size set of nodes once per layer, shared by all nodes in that layer.
- Therefore the cost grows only linearly with depth instead of exponentially.

Degrees: $d(A)=3$, $d(B)=3$, $d(C)=1$, $d(D)=3$, $d(E)=2$

Divide each row by its degree. $\tilde{A} = [0, 0.3333, 0.3333, 0.3333, 0]; [0.3333, 0, 0, 0.3333, 0.3333]; [1, 0, 0, 0, 0]; [0.3333, 0.3333, 0, 0, 0.3333]; [0, 0.5, 0, 0.5, 0]$ (each row sums to 1)

$q(u)$ uses the squared norm of column u of $\tilde{A}$, not of a row.

Column A: squared sum = 1.2222

Column B: squared sum = 0.4722

Column C: squared sum = 0.1111

Column D: squared sum = 0.4722

Column E: squared sum = 0.2222

Divide by the total 2.5 to turn them into probabilities.

b) $q(A) = 0.4889$, $q(B) = 0.1889$, $q(C) = 0.0444$, $q(D) = 0.1889$, $q(E) = 0.0889$

For (c) look only at the sampled nodes: read $\tilde{A}(A, u)$ for each.

$\tilde{A}(A, D) h(D) = 0.3333 \times 1 = 0.3333$

$\tilde{A}(A, E) h(E) = 0 \times -1 = 0$

Add them and divide by the number sampled (2), not by the degree: $(0.3333 + 0)/2 = 0.1667$

No negative value, so ReLU changes nothing.

c) $h_1(A) = 0.1667$

3. (ASGCN) Consider the undirected graph G with adjacency matrix A and node feature vector $h^{(0)}$ below. ASGCN rewrites a GCN layer as an expectation:

$$A = \begin{bmatrix} 0 & 1 & 1 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 1 & 1 \\ 0 & 1 & 1 & 0 & 1 \\ 0 & 0 & 1 & 1 & 0 \end{bmatrix}$$

$$h^{(0)} = \begin{bmatrix} 4 \\ -1 \\ 2 \\ 2 \\ 1 \end{bmatrix}$$

$$h^{(l+1)}(v_i) = \sigma(W^{(l)} N(v_i) E_{p(u_j | v_i)}[h^{(l)}(u_j)])$$

where $\tilde{A} = D^{-1}A$, $N(v_i) = \sum_j \tilde{A}_{ij}$, $p(u_j | v_i) = \tilde{A}_{ij} / N(v_i)$, $W^{(l)} = 1$ and $\text{ReLU}(x) = \max(0, x)$.

- What does the proposal distribution depend on in ASGCN that it does not in FastGCN?
- Calculate $N(v_i)$ and $p(u_j | v_i)$ for the parent node $v_i = C$.
- Calculate $E_{p(u_j | C)}[h^{(0)}(u_j)]$ and the resulting $h^{(1)}(C)$.

Solution

ENG)

- Write these three sentences.
 - In FastGCN the sampling distribution is fixed in advance and does not depend on the parent nodes.
 - In ASGCN the proposal distribution is conditioned on the nodes of the upper layer.
 - This makes almost every sampled node useful and reduces the variance of the estimate.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=3$, $d(D)=3$, $d(E)=2$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0.5, 0, 0]; [0.5, 0, 0, 0.5, 0]; [0.3333, 0, 0, 0.3333, 0.3333]; [0, 0.3333, 0.3333, 0, 0.3333]; [0, 0, 0.5, 0.5, 0]$ (each row sums to 1)

The parent node is C, so only its row of $\tilde{A}$ matters.

$N(C) = \text{sum of that row} = 1$. With row normalisation it is 1.

$$p(A|C) = \tilde{A}_{CA} / N = 0.3333$$

$$p(D|C) = \tilde{A}_{CD} / N = 0.3333$$

$$p(E|C) = \tilde{A}_{CE} / N = 0.3333$$

$$\text{b) } N(C) = 1, p(A|C) = 0.3333, p(D|C) = 0.3333, p(E|C) = 0.3333$$

The expectation is the sum of probability times value: $0.3333 \times 4 + 0.3333 \times 2 + 0.3333 \times 1 = 2.3333$

$N = 1$ and $W = 1$, so the value is unchanged. It equals the plain neighbour mean.

No negative value, so ReLU changes nothing.

$$\text{c) } E = 2.3333, h_1(C) = 2.3333$$

4. (APPNP) Given a graph with an adjacency matrix A and a prediction matrix H produced by an MLP as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 \\ 1 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

$$H = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$$

APPNP propagates the predictions with a teleport probability α :

$$Z^{(0)} = H, \quad Z^{(k+1)} = (1 - \alpha) \tilde{A} Z^{(k)} + \alpha H$$

where $\tilde{A} = D^{-1}A$ and $\alpha = 0.5$. No nonlinearity is applied between steps.

- Explain the teleport mechanism and why it prevents over-smoothing.
- Calculate $\tilde{A}$ and $Z^{(1)}$.
- Calculate $Z^{(2)}$.

Solution

ENG)

a) Write these three sentences.

- At every propagation step the walk returns to the root node with probability alpha. This is the teleport mechanism.
- Because a fixed fraction of the mass is pulled back at every step, the distribution stays local.
- So we can use many propagation steps without over-smoothing, and propagation has no extra parameters.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=3$, $d(D)=1$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0.5, 0]; [0.5, 0, 0.5, 0]; [0.3333, 0.3333, 0, 0.3333]; [0, 0, 1, 0]$
(each row sums to 1)

b) $\tilde{A} = [0, 0.5, 0.5, 0]; [0.5, 0, 0.5, 0]; [0.3333, 0.3333, 0, 0.3333]; [0, 0, 1, 0]$

Start from $Z_0 = H = [0, 0, 0, 1]$

Propagate one step: $\tilde{A} Z_0 = [0, 0, 0.3333, 0]$

Multiply by $(1-\alpha) = 0.5$: $[0, 0, 0.1667, 0]$

Add $\alpha H = [0, 0, 0, 0.5]$. This is the teleport term.

$$Z_1 = [0, 0, 0.1667, 0.5]$$

One more step: $\tilde{A} Z_1 = [0.0833, 0.0833, 0.1667, 0.1667]$

Multiply by $(1-\alpha)$ and add αH **again**. It is added at every step.

c) $Z_2 = [0.0417, 0.0417, 0.0833, 0.5833]$

5. (GPR-GNN) Given a graph with an adjacency matrix A and hidden features $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 3 \\ 0 \\ -1 \\ -2 \end{bmatrix}$$

GPR-GNN separates feature extraction from topological propagation:

$$H^{(k)} = \tilde{A} H^{(k-1)}, \quad Z = \sum_{k=0}^K \gamma_k H^{(k)}$$

where $\tilde{A} = D^{-1}A$, $K = 2$, $\gamma_0 = 0.7$, $\gamma_1 = 0.3$, $\gamma_2 = -0.2$.

- What can GPR-GNN do that APPNP cannot, and why do negative γ_k matter?
- Calculate $H^{(1)}$ and $H^{(2)}$.
- Calculate Z .

Solution

ENG)

- Write these three sentences.
 - APPNP fixes the coefficients to alpha times (1-alpha) to the power k, which are always positive and decreasing.
 - GPR-GNN learns the weights, and they are allowed to be negative.
 - A negative weight subtracts that hop, so the model can act as a high-pass filter on heterophilic graphs.

Degrees: $d(A)=2$, $d(B)=1$, $d(C)=2$, $d(D)=3$

Divide each row by its degree. $\tilde{A} = [0, 0, 0.5, 0.5]; [0, 0, 0, 1]; [0.5, 0, 0, 0.5]; [0.3333, 0.3333, 0.3333, 0]$
(each row sums to 1)

H_0 stays as it is: $[3, 0, -1, -2]$

$H_1 = \tilde{A} H_0 = [-1.5, -2, 0.5, 0.6667]$

$H_2 = \tilde{A} H_1 = [0.5833, 0.6667, -0.4167, -1]$

b) $H_1 = [-1.5, -2, 0.5, 0.6667]$, $H_2 = [0.5833, 0.6667, -0.4167, -1]$

Now take the weighted sum with gamma. There is no ReLU during propagation.

$\gamma_0 H_0 = 0.7 \times [3, 0, -1, -2] = [2.1, 0, -0.7, -1.4]$

$\gamma_1 H_1 = 0.3 \times [-1.5, -2, 0.5, 0.6667] = [-0.45, -0.6, 0.15, 0.2]$

$\gamma_2 H_2 = -0.2 \times [0.5833, 0.6667, -0.4167, -1] = [-0.1167, -0.1333, 0.0833, 0.2]$

γ_2 is negative, so that hop is subtracted. Dropping the sign gives a wrong answer.

Add the three vectors element-wise.

c) $Z = [1.5333, -0.7333, -0.4667, -1]$

6. (MixHop) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 0 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 2 \\ 3 \\ 4 \\ -1 \end{bmatrix}$$

A MixHop layer concatenates several powers of the adjacency matrix:

$$H^{(i+1)} = \parallel_{j \in P} \sigma(\hat{A}^j H^{(i)} W_j^{(i)})$$

where $P = \{0, 1, 2\}$, $\hat{A} = D^{-1}A$, $W_j^{(i)} = 1$ and $\text{ReLU}(x) = \max(0, x)$.

- State Theorem 1 of MixHop about the two-hop Delta Operator.
- Calculate $\hat{A}^0 H^{(0)}$, $\hat{A}^1 H^{(0)}$ and $\hat{A}^2 H^{(0)}$.
- Calculate $H^{(1)}$ and the Delta Operator $\sigma(\hat{A} H^{(0)}) - \sigma(\hat{A}^2 H^{(0)})$.

Solution

ENG)

a) Write these three sentences.

- Theorem 1 says that a vanilla GCN cannot represent the two-hop Delta Operator, even with infinitely many layers.
- Stacking GCN layers collapses to $\hat{A}$ to the power l times W , so the intermediate hops are never kept separately.
- MixHop computes each power separately and concatenates them, so a later layer can subtract one from another.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=1$, $d(D)=3$

Divide each row by its degree. $\text{Atilde} = [0, 0.5, 0, 0.5]; [0.5, 0, 0, 0.5]; [0, 0, 0, 1]; [0.3333, 0.3333, 0.3333, 0]$ (each row sums to 1)

$\hat{A}^0$ is the identity, so it is just H^0 : $[2, 3, 4, -1]$

$\hat{A}^1 H^0 = [1, 0.5, -1, 3]$ (one hop, the neighbour mean)

$\hat{A}^2 H^0 = \text{Atilde}(\text{Atilde } H^0) = [1.75, 2, 3, 0.1667]$ (two hops)

ReLU is applied to each block separately, not after concatenation.

After ReLU: $[2, 3, 4, 0] \parallel [1, 0.5, 0, 3] \parallel [1.75, 2, 3, 0.1667]$

The delta operator subtracts the two-hop block from the one-hop block.

b) $A^0 H = [2, 3, 4, -1]$, $A^1 H = [1, 0.5, -1, 3]$, $A^2 H = [1.75, 2, 3, 0.1667]$

c) $H^1 =$ concatenation of the three blocks: $[2, 3, 4, 0] \parallel [1, 0.5, 0, 3] \parallel [1.75, 2, 3, 0.1667]$

Delta = $[-0.75, -1.5, -3, 2.8333]$

7. (**PairNorm**) After one GCN layer the node features of a four-node graph are

$$X = \begin{bmatrix} 0 & 4 \\ 1 & 3 \\ 1 & -2 \\ -2 & 1 \end{bmatrix}$$

PairNorm normalizes them with a centering step and a scaling step:

$$x_i^c = x_i - (1/n) \sum_i x_i \quad (\text{Center})$$

$$x_i = s \cdot x_i^c / \text{sqrt}((1/n) \sum_i \|x_i^c\|_2^2) \quad (\text{Scale})$$

with $s = 1$ and $n = 4$.

- Explain node-wise and feature-wise over-smoothing, and what PairNorm keeps constant.
- Calculate the centered features.
- Calculate the scaling denominator and the final x_i .

Solution

ENG)

- Write these three sentences.
 - Node-wise over-smoothing means the features of connected nodes become too similar.
 - Feature-wise over-smoothing means all node features become indistinguishable across the whole graph.
 - PairNorm keeps the total pairwise squared distance between node representations constant across layers.

First take the mean over the nodes, separately for each feature column.

Column 1: $(0 + 1 + 1 + -2)/4 = 0$

Column 2: $(4 + 3 + -2 + 1)/4 = 1.5$

Subtract the mean from every row. Check: each column now sums to zero.

Node 1: squared norm = 6.25

Node 2: squared norm = 3.25

Node 3: squared norm = 13.25

Node 4: squared norm = 4.25

Mean = $(6.25 + 3.25 + 13.25 + 4.25)/4 = 6.75$

Take the square root: $\text{sqrt}(6.75) = 2.5981$. Divide every row by this single number.

b) Mean = $[0, 1.5]$, centered = $[0, 2.5]; [1, 1.5]; [1, -3.5]; [-2, -0.5]$

c) Scaling denominator = 2.5981, $x = [0, 0.9623]; [0.3849, 0.5774]; [0.3849, -1.3472]; [-0.7698, -0.1925]$

8. (DropEdge) Consider the complete graph K_4 on nodes A, B, C and D with the adjacency matrix A and initial node feature matrix $H^{(0)}$ below.

$$A = \begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 0 \\ -2 \\ 2 \\ 2 \end{bmatrix}$$

DropEdge keeps each undirected edge independently with probability $1 - p$ at every epoch:

$$H^{(1)} = \sigma((D^{(t)})^{-1} A^{(t)} H^{(0)})$$

where $A^{(t)}$ and $D^{(t)}$ are rebuilt from the surviving edges only, $p = 1/3$, and $\text{ReLU}(x) = \max(0, x)$.

- Which two problems does DropEdge target, and how does it differ from Dropout?
- Calculate the expected number of surviving edges, and $H^{(1)}$ on the full graph.
- Suppose the edges (C, D) and (A, C) are dropped. Recalculate $H^{(1)}$.

Solution

ENG)

- Write these three sentences.
 - DropEdge targets over-fitting and over-smoothing.
 - A different random subgraph is used at every epoch, which works as data augmentation on the graph.
 - Dropout zeroes entries of the feature tensor, while DropEdge removes edges of the adjacency matrix.

K_4 has 6 edges. Each survives with probability $1 - 1/3 = 2/3$, so the expectation is 4.

First the full graph: every node has degree 3.

A: neighbours B, C, D. $(-2 + 2 + 2)/3 = 0.6667$

B: neighbours A, C, D. $(0 + 2 + 2)/3 = 1.3333$

C: neighbours A, B, D. $(0 + -2 + 2)/3 = 0$

D: neighbours A, B, C. $(0 + -2 + 2)/3 = 0$

Now drop the two given edges and **recount the degrees**. This is where most mistakes happen.

Degrees after dropping: $d(A)=2$, $d(B)=3$, $d(C)=1$, $d(D)=2$

A: neighbours B, D. $(-2 + 2)/2 = 0$

B: neighbours A, C, D. $(0 + 2 + 2)/3 = 1.3333$

C: neighbours B. $(-2)/1 = -2$

D: neighbours A, B. $(0 + -2)/2 = -1$

b) Expected number of surviving edges = 4, $H_1 = [0.6667, 1.3333, 0, 0]$

c) $H_1 = [0, 1.3333, 0, 0]$

9. (Node2vec) Consider an undirected graph G given in the following figure. A random walk has just moved from node A to node B , with $p = 2$ and $q = 0.5$.




Figure 1. Graph G

The unnormalized transition probability from B to a next node x is $\pi_{Bx} = \alpha_{pq}(A, x)$, where $\alpha = 1/p$ if $d(A, x) = 0$, $\alpha = 1$ if $d(A, x) = 1$, and $\alpha = 1/q$ if $d(A, x) = 2$.

- Explain BFS and DFS, and explain which one a small q makes the walk behave like.
- List all possible next nodes from node B and calculate π_{Bx} .
- Calculate the normalized transition probabilities.

Solution

ENG)

- Write these three sentences.
 - BFS explores the immediate neighbourhood first, so the embedding captures the structural role of a node.
 - DFS moves far away from the source, so the embedding captures community membership.
 - A small q makes $1/q$ large, so the walk moves outward and behaves like DFS.

The previous node is A and the walk is now at B . All distances are measured from A .

A : $x = A$ is the previous node, $d = 0$, $\alpha = 1/p = 0.5$

C : x is two hops from A , $d = 2$, $\alpha = 1/q = 2$

E : x is also a neighbour of A , $d = 1$, $\alpha = 1 = 1$

Sum them: $0.5 + 2 + 1 = 3.5$

b) $\pi(B \rightarrow A) = 0.5$, $\pi(B \rightarrow C) = 2$, $\pi(B \rightarrow E) = 1$

c) $P(A) = 0.1429$, $P(C) = 0.5714$, $P(E) = 0.2857$

10. (LINE) Consider an undirected graph G with node embeddings

$$u_A = 1, u_B = 0.5, u_C = -0.5, u_D = -1$$

and context embeddings

$$u'_A = -0.5, u'_B = 1, u'_C = 0.5, u'_D = -1.$$

The second-order proximity is $p_2(j|x) = \exp(u'_j u_x) / \sum_k \exp(u'_k u_x)$. Given $e^{0.25} = 1.28$, $e^{0.5} = 1.65$, $e^{-0.25} = 0.78$, $e^{-0.5} = 0.61$.

- What is second-order proximity, and how does it differ from first-order proximity?
- Calculate $p_2(A|C)$, $p_2(B|C)$, $p_2(C|C)$ and $p_2(D|C)$.

Solution

ENG)

- Write these three sentences.

- First-order proximity is the similarity between two nodes that are directly connected by an edge.
- Second-order proximity is the similarity between two nodes that share many neighbours, even without an edge.
- So LINE gives each node two vectors: its own embedding and its context embedding.

The base node is C with $u(C) = -0.5$. The other side uses the context vector u' .

$$A: u'(A) u(C) = -0.5 \times -0.5 = 0.25, \exp = 1.28$$

$$B: u'(B) u(C) = 1 \times -0.5 = -0.5, \exp = 0.61$$

$$C: u'(C) u(C) = 0.5 \times -0.5 = -0.25, \exp = 0.78$$

$$D: u'(D) u(C) = -1 \times -0.5 = 0.5, \exp = 1.65$$

The denominator is the sum of all four: $1.28 + 0.61 + 0.78 + 1.65 = 4.32$

$$b) p_2(A|C) = 0.2963, p_2(B|C) = 0.1412, p_2(C|C) = 0.1806, p_2(D|C) = 0.3819$$

11. (GCN) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 & 1 \\ 0 & 0 & 1 & 0 & 1 \\ 0 & 1 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} -1 \\ -2 \\ 2 \\ -2 \\ 4 \end{bmatrix}$$

Assume that the hidden layer of a GCN model of all nodes at layer (k) can be calculated as:

$$H^{(k)} = \sigma(\tilde{A} \cdot H^{(k-1)}),$$

where $\tilde{A} = D^{-1}A$ and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- Explain how information is propagated in a GCN layer.
- Calculate $\tilde{A}$.
- Calculate the output of the GCN model at layer $k = 1$.

Solution

ENG)

a) Write these three sentences.

- In one layer each node aggregates the features of its direct neighbours and takes the average.
- The row-normalized matrix $\tilde{A}$ equals $D^{-1}A$, which performs this averaging.
- After k layers a node has received information from all nodes within k hops.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=3$, $d(D)=2$, $d(E)=3$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0.5, 0, 0]; [0.5, 0, 0, 0, 0.5]; [0.3333, 0, 0, 0.3333, 0.3333]; [0, 0, 0.5, 0, 0.5]; [0, 0.3333, 0.3333, 0.3333, 0]$ (each row sums to 1)

b) $\tilde{A} = [0, 0.5, 0.5, 0, 0]; [0.5, 0, 0, 0, 0.5]; [0.3333, 0, 0, 0.3333, 0.3333]; [0, 0, 0.5, 0, 0.5]; [0, 0.3333, 0.3333, 0.3333, 0]$

Multiply $\tilde{A}$ by H^0 . Each entry is the mean of that node's neighbours.

A: neighbours B, C. $(-2 + 2)/2 = 0$

B: neighbours A, E. $(-1 + 4)/2 = 1.5$

C: neighbours A, D, E. $(-1 + -2 + 4)/3 = 0.3333$

D: neighbours C, E. $(2 + 4)/2 = 3$

E: neighbours B, C, D. $(-2 + 2 + -2)/3 = -0.6667$

ReLU clips the negatives: $-0.6667 \rightarrow 0$

c) $H_1 = [0, 1.5, 0.3333, 3, 0]$

12. (GraphSAGE) Consider an undirected graph G of five nodes given in the following figure. Each node has initial features that are the numbers standing next to it. According to the GraphSAGE model with an AGGREGATE that is a MEAN function, the feature of a node i at layer k can be updated as:

$$h_{N(i)}^{(k)} = \text{MEAN}(\{ h_u^{(k-1)}, \forall u \in N(i) \})$$

$$h_i^{(k)} = \text{ReLU}(h_i^{(k-1)} \parallel h_{N(i)}^{(k)})$$

where $\parallel$ is a concatenation and $\text{ReLU}(x) = \max(0, x)$.




Figure. Graph G

- What is the main idea behind GraphSAGE, and how does it differ from traditional GCNs?
- Calculate the feature of each node at $k = 1$.
- Calculate a graph-level embedding h_G by using a 'Mean' global pooling when $k = 1$.

Solution

ENG)

a) Write these three sentences.

- GraphSAGE learns an aggregator function instead of a fixed embedding for each node.
- It samples a fixed number of neighbours, so it can be trained in mini-batches on very large graphs.
- A GCN is transductive, but GraphSAGE is inductive, so it can embed a node that was never seen in training.

First compute the neighbour mean. The node itself is not included.

A: neighbours B, C, D. $(3 + -1 + -2)/3 = 0$

B: neighbours A, C, E. $(-2 + -1 + -2)/3 = -1.6667$

C: neighbours A, B. $(-2 + 3)/2 = 0.5$

D: neighbours A, E. $(-2 + -2)/2 = -2$

E: neighbours B, D. $(3 + -2)/2 = 0.5$

Concatenate the node's own value with the neighbour mean. This is not an addition.

ReLU clips the negatives: $-2 \rightarrow 0$, $-1.6667 \rightarrow 0$, $-1 \rightarrow 0$, $-2 \rightarrow 0$, $-2 \rightarrow 0$, $-2 \rightarrow 0$

The graph embedding is the mean per position: first $-4 / 5$, second $-2.6667 / 5$

b) $h(A) = [0, 0]$, $h(B) = [3, 0]$, $h(C) = [0, 0.5]$, $h(D) = [0, 0]$, $h(E) = [0, 0.5]$

c) $h_G = [0.6, 0.2]$

13. (ClusterGCN) Consider an undirected graph G of six nodes given in the following figure. The graph G contains two clusters G_1 and G_2 . Each node has initial features that are the numbers standing next to it. According to the ClusterGCN model, the feature of a node i at layer k can be updated as:

$$h_{N(i)}^{(k)} = \text{MEAN}(\{ h_u^{(k-1)}, \forall u \in N(i), G_u = G_i \})$$

$$h_i^{(k)} = \text{ReLU}(h_i^{(k-1)} \parallel h_{N(i)}^{(k)})$$

where $\parallel$ is a concatenation and $\text{ReLU}(x) = \max(0, x)$.




Figure. Graph G

- What is mini-batch training, and why does ClusterGCN need it?
- Calculate the output representations of all nodes at layer $k = 1$.

Solution

ENG)

- Write these three sentences.

- Mini-batch training updates the parameters using a small subset of the data at a time.
- A large graph does not fit in memory, so ClusterGCN partitions the graph into densely connected clusters.
- A good partition keeps most edges inside a batch, so little information is lost at the boundary.

Edges that cross a group are treated as absent. Cross them out first.

A: neighbours B. $(4)/1 = 4$

B: neighbours A. $(3)/1 = 3$

C: no neighbour inside the group, so the mean is 0.

D: no neighbour inside the group, so the mean is 0.

E: neighbours F. $(-2)/1 = -2$

F: neighbours E. $(1)/1 = 1$

Concatenate the neighbour mean after the node's own value, then apply ReLU to both entries.

- $h(A) = [3, 4]$, $h(B) = [4, 3]$, $h(C) = [0, 0]$, $h(D) = [2, 0]$, $h(E) = [1, 0]$, $h(F) = [0, 1]$

14. (GraphSAINT) Consider an undirected graph G of six nodes given in the following figure. The graph G has a sampled subgraph G_1 . Each node has initial features that are the numbers standing next to it. According to the GraphSAINT model, the feature of a node i at layer k can be updated as:

$$h_{N(i)}^{(k)} = \text{MEAN}(\{ h_u^{(k-1)}, \forall u \in N(i), G_u = G_i \})$$

$$h_i^{(k)} = \text{ReLU}(h_i^{(k-1)} \parallel h_{N(i)}^{(k)})$$

where $\parallel$ is a concatenation and $\text{ReLU}(x) = \max(0, x)$.




Figure. Graph G

- How does GraphSAINT's sampling differ from node-wise or layer-wise sampling?
- Calculate the output representations of nodes A, B, C and D at layer $k = 1$.

Solution

ENG)

- Write these three sentences.
 - GraphSAGE and FastGCN sample while building the layers, so sampling is repeated at every layer.
 - GraphSAINT samples a subgraph first and then builds the full GNN on that subgraph.
 - So the cost does not grow with depth, and normalization coefficients correct the sampling bias.

Edges that leave the sampled subgraph are discarded.

A: neighbours B, D. $(4 + -1)/2 = 1.5$

B: neighbours A, C. $(4 + 2)/2 = 3$

C: neighbours B. $(4)/1 = 4$

D: neighbours A. $(4)/1 = 4$

Concatenate the neighbour mean after the node's own value, then apply ReLU.

b) $h(A) = [4, 1.5]$, $h(B) = [4, 3]$, $h(C) = [2, 4]$, $h(D) = [0, 4]$

15. (JK Network) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 0 & 1 & 1 \\ 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 1 & 1 & 0 & 0 & 1 \\ 1 & 0 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 3 \\ 0 \\ -3 \\ -1 \\ -1 \end{bmatrix}$$

Assume that the output of a JK network model of all nodes at layer (k) can be calculated as:

$$H^{(k)} = \max(\sigma(\tilde{A} H^{(0)}), \sigma(\tilde{A} H^{(1)}), \dots, \sigma(\tilde{A} H^{(k-1)}))$$

where $\tilde{A} = D^{-1}A$ and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- Explain how Jumping Knowledge Networks help mitigate the over-smoothing problem.
- Calculate $\tilde{A}$ and $H^{(1)}$.
- Calculate the output representations at layer $k = 2$.

Solution

ENG)

a) Write these three sentences.

- In a deep GCN every node repeatedly averages its neighbours, so all representations become similar. This is over-smoothing.
- JKNet keeps the representation of every layer and combines them at the end by max, concatenation or an LSTM.
- The best depth is different for each node, so each node can select the depth that suits it.

Degrees: $d(A)=3$, $d(B)=2$, $d(C)=1$, $d(D)=3$, $d(E)=3$

Divide each row by its degree. $\text{Atilde} = [0, 0.3333, 0, 0.3333, 0.3333]; [0.5, 0, 0, 0.5, 0]; [0, 0, 0, 0, 1]; [0.3333, 0.3333, 0, 0, 0.3333]; [0.3333, 0, 0.3333, 0.3333, 0]$ (each row sums to 1)

b) $\text{Atilde} = [0, 0.3333, 0, 0.3333, 0.3333]; [0.5, 0, 0, 0.5, 0]; [0, 0, 0, 0, 1]; [0.3333, 0.3333, 0, 0, 0.3333]; [0.3333, 0, 0.3333, 0.3333, 0]$

First build $H1 = \text{ReLU}(\text{Atilde } H0)$

A: neighbours B, D, E. $(0 + -1 + -1)/3 = -0.6667$

B: neighbours A, D. $(3 + -1)/2 = 1$

C: neighbours E. $(-1)/1 = -1$

D: neighbours A, B, E. $(3 + 0 + -1)/3 = 0.6667$

E: neighbours A, C, D. $(3 + -3 + -1)/3 = -0.3333$

ReLU clips the negatives: $-0.6667 \rightarrow 0$, $-1 \rightarrow 0$, $-0.3333 \rightarrow 0$

$H1 = [0, 1, 0, 0.6667, 0]$

For $k = 2$ there are two terms. First: $\text{ReLU}(\text{Atilde } H0) = [0, 1, 0, 0.6667, 0]$

Second: $\text{ReLU}(\text{Atilde } H1) = [0.5556, 0.3333, 0, 0.3333, 0.2222]$

The max is taken element-wise, not by choosing one whole vector.

c) $H_2 = [0.5556, 1, 0, 0.6667, 0.2222]$

16. (GCNII) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 0 & 0 & 1 & 1 \\ 0 & 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 \\ 1 & 1 & 0 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 0 \\ -3 \\ -1 \\ 1 \\ -1 \end{bmatrix}$$

Assume that the output of a GCNII model of all nodes at layer (k) can be calculated as:

$$H^{(k)} = \sigma[((1 - \beta) I_n) \cdot ((1 - \alpha) \tilde{A} \cdot H^{(k-1)} + \alpha H^{(0)})]$$

where $\tilde{A} = D^{-1}A$, I_n is the identity matrix, $\alpha = 0.2$, $\beta = 0.4$, and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- What are the two major components introduced in GCNII to overcome the depth limitation?
- Calculate $\tilde{A}$.
- Calculate the output representations at layer $k = 1$.

Solution

ENG)

- Write these three sentences.
 - The two components are the initial residual connection and the identity mapping.
 - The initial residual connection adds a fraction alpha of the input H zero back at every layer.
 - The identity mapping keeps the weight matrix close to the identity, so a deep model is not destroyed.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=2$, $d(D)=3$, $d(E)=3$

Divide each row by its degree. $\tilde{A} = [0, 0, 0, 0.5, 0.5]; [0, 0, 0.5, 0, 0.5]; [0, 0.5, 0, 0.5, 0]; [0.3333, 0, 0.3333, 0, 0.3333]; [0.3333, 0.3333, 0, 0.3333, 0]$ (each row sums to 1)

b) $\tilde{A} = [0, 0, 0, 0.5, 0.5]; [0, 0, 0.5, 0, 0.5]; [0, 0.5, 0, 0.5, 0]; [0.3333, 0, 0.3333, 0, 0.3333]; [0.3333, 0.3333, 0, 0.3333, 0]$

Start with the inner bracket. $\tilde{A} H^0 = [0, -1, -1, -0.6667, -0.6667]$

Multiply by $(1-\alpha) = 0.8$: $[0, -0.8, -0.8, -0.5333, -0.5333]$

Add $\alpha H^0 = 0.2 \times [0, -3, -1, 1, -1] = [0, -0.6, -0.2, 0.2, -0.2]$

Inner bracket = $[0, -1.4, -1, -0.3333, -0.7333]$

Multiply by the outer $(1-\beta) = 0.6$. Dropping this factor scales the answer.

ReLU clips the negatives: $-0.84 \rightarrow 0$, $-0.6 \rightarrow 0$, $-0.2 \rightarrow 0$, $-0.44 \rightarrow 0$

c) $H^1 = [0, 0, 0, 0, 0]$

17. (DeepGCN) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 0 & 1 \\ 1 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} -2 \\ 0 \\ 3 \\ 1 \end{bmatrix}$$

Assume that the output of a DeepGCN model with a residual connection at layer (k) is:

$$H^{(k)} = \sigma(\tilde{A} \cdot H^{(k-1)}) + H^{(k-1)}$$

where $\tilde{A} = D^{-1}A$ and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- Explain why a residual connection allows a GCN to be made deeper.
- Calculate the output representations at layer $k = 1$ and $k = 2$.
- Compare $H^{(1)}$ with the output of a plain GCN without the residual connection.

Solution

ENG)

- Write these three sentences.
 - Deep GCNs suffer from vanishing gradients and over-smoothing, so adding layers makes the model worse.
 - A residual connection passes the input forward unchanged, so each layer only has to learn the change.
 - If no change is useful the layer can output zero, and gradients also flow back through the skip path.

Degrees: $d(A)=2$, $d(B)=3$, $d(C)=2$, $d(D)=3$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0, 0.5]; [0.3333, 0, 0.3333, 0.3333]; [0, 0.5, 0, 0.5]; [0.3333, 0.3333, 0.3333, 0]$ (each row sums to 1)

$\tilde{A} H^0 = [0.5, 0.6667, 0.5, 0.3333]$

Apply ReLU first: $[0.5, 0.6667, 0.5, 0.3333]$

Then add H^0 . The addition sits outside the ReLU, so negatives can survive.

For $k = 2$, feed H^1 back in and add H^1 , not H^0 .

$\tilde{A} H^1 = [1, 1.1111, 1, 0.8889]$, after ReLU $[1, 1.1111, 1, 0.8889]$

b) $H^1 = [-1.5, 0.6667, 3.5, 1.3333]$

$H^2 = [-0.5, 1.7778, 4.5, 2.2222]$

c) Plain GCN without residual: $H^1 = [0.5, 0.6667, 0.5, 0.3333]$ (the node's own value and negatives disappear)

3회차

1. (Relational Classification) Consider an undirected graph G of five nodes given in the following figure. Nodes 1 and 5 are labelled ($P(Y_1) = 1, P(Y_5) = 0$) and the others are unlabelled and initialized to 0.5. The relational classifier updates every unlabelled node as:

$$P(Y_v = c) = (\sum_{(v,u) \in E} A_{v,u} P(Y_u = c)) / (\sum_{(v,u) \in E} A_{v,u})$$




Figure. Graph G

- Explain the key assumption behind relational classification, and state one limitation.
- Calculate $P(Y_v)$ for every unlabelled node after the 1st iteration.
- Calculate $P(Y_v)$ after the 2nd iteration, using the previous iteration's values only.

Solution

ENG)

- Write these three sentences.
 - Relational classification assumes homophily: connected nodes tend to share the same label.
 - It repeatedly sets each unlabelled node to the weighted average of its neighbours' class probabilities.
 - Its limitation is that it uses only the graph structure and ignores node features.

Write the initial values: node 1 is 1, node 5 is 0, the rest start at 0.5.

First iteration. Only the unlabelled nodes are replaced by the mean of their neighbours.

Node 1 is labelled, so it never changes. $P = 1$

Node 2: neighbours 1, 4, 5. $(1 + 0.5 + 0)/3 = 0.5$

Node 3: neighbours 1, 5. $(1 + 0)/2 = 0.5$

Node 4: neighbours 1, 2. $(1 + 0.5)/2 = 0.75$

Node 5 is labelled, so it never changes. $P = 0$

b) $P(Y_1) = 1, P(Y_2) = 0.5, P(Y_3) = 0.5, P(Y_4) = 0.75, P(Y_5) = 0$

Second iteration. Use only the values from the first iteration, not the ones just computed.

Node 1 is labelled, so it never changes. $P = 1$

Node 2: neighbours 1, 4, 5. $(1 + 0.75 + 0)/3 = 0.5833$

Node 3: neighbours 1, 5. $(1 + 0)/2 = 0.5$

Node 4: neighbours 1, 2. $(1 + 0.5)/2 = 0.75$

Node 5 is labelled, so it never changes. $P = 0$

c) $P(Y_1) = 1, P(Y_2) = 0.5833, P(Y_3) = 0.5, P(Y_4) = 0.75, P(Y_5) = 0$

2. (FastGCN) Consider the undirected graph G with adjacency matrix A and node feature vector $h^{(0)}$ below. FastGCN replaces the full neighbourhood sum of a GCN layer by a layer-wise Monte-Carlo estimate:

$$A = \begin{bmatrix} 0 & 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 1 & 1 \\ 0 & 1 & 1 & 0 & 1 \\ 0 & 1 & 1 & 1 & 0 \end{bmatrix}$$

$$h^{(0)} = \begin{bmatrix} 1 \\ 0 \\ -2 \\ -1 \\ 2 \end{bmatrix}$$

$$h^{(l+1)}(v) = \sigma \left(\frac{1}{t_l} \sum_{j=1}^{t_l} \tilde{A}(v, u_j) h^{(l)}(u_j) \right)$$

where $\tilde{A} = D^{-1}A$ and the t_l nodes are drawn once per layer from $q(u) \propto \|\tilde{A}(:, u)\|^2$, with $\text{ReLU}(x) = \max(0, x)$.

- How does layer-wise sampling differ from the node-wise sampling of GraphSAGE?
- Calculate the importance distribution $q(u)$ over the five nodes.
- Suppose the sampled set is $\{B, C\}$ ($t_l = 2$). Calculate the estimate of $h^{(1)}(A)$.

Solution

ENG)

a) Write these three sentences.

- In node-wise sampling, each node samples its own neighbours, so the cost grows exponentially with depth.
- FastGCN samples a fixed-size set of nodes once per layer, shared by all nodes in that layer.
- Therefore the cost grows only linearly with depth instead of exponentially.

Degrees: $d(A)=1$, $d(B)=3$, $d(C)=2$, $d(D)=3$, $d(E)=3$

Divide each row by its degree. $\text{Atilde} = [0, 1, 0, 0, 0]; [0.3333, 0, 0, 0.3333, 0.3333]; [0, 0, 0, 0.5, 0.5]; [0, 0.3333, 0.3333, 0, 0.3333]; [0, 0.3333, 0.3333, 0.3333, 0]$ (each row sums to 1)

$q(u)$ uses the squared norm of column u of Atilde , not of a row.

Column A: squared sum = 0.1111

Column B: squared sum = 1.2222

Column C: squared sum = 0.2222

Column D: squared sum = 0.4722

Column E: squared sum = 0.4722

Divide by the total 2.5 to turn them into probabilities.

b) $q(A) = 0.0444$, $q(B) = 0.4889$, $q(C) = 0.0889$, $q(D) = 0.1889$, $q(E) = 0.1889$

For (c) look only at the sampled nodes: read $\text{Atilde}(A, u)$ for each.

$\text{Atilde}(A, B) h(B) = 1 \times 0 = 0$

$\text{Atilde}(A, C) h(C) = 0 \times -2 = 0$

Add them and divide by the number sampled (2), not by the degree: $(0 + 0)/2 = 0$

No negative value, so ReLU changes nothing.

c) $h_1(A) = 0$

3. (ASGCN) Consider the undirected graph G with adjacency matrix A and node feature vector $h^{(0)}$ below. ASGCN rewrites a GCN layer as an expectation:

$$A = \begin{bmatrix} 0 & 0 & 1 & 0 & 1 \\ 0 & 0 & 0 & 1 & 1 \\ 1 & 0 & 0 & 1 & 1 \\ 0 & 1 & 1 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \end{bmatrix}$$

$$h^{(0)} = \begin{bmatrix} -1 \\ -1 \\ 4 \\ 1 \\ -2 \end{bmatrix}$$

$$h^{(l+1)}(v_i) = \sigma(W^{(l)} N(v_i) E_{p(u_j | v_i)}[h^{(l)}(u_j)])$$

where $\tilde{A} = D^{-1}A$, $N(v_i) = \sum_j \tilde{A}_{ij}$, $p(u_j | v_i) = \tilde{A}_{ij} / N(v_i)$, $W^{(l)} = 1$ and $\text{ReLU}(x) = \max(0, x)$.

- What does the proposal distribution depend on in ASGCN that it does not in FastGCN?
- Calculate $N(v_i)$ and $p(u_j | v_i)$ for the parent node $v_i = C$.
- Calculate $E_{p(u_j | C)}[h^{(0)}(u_j)]$ and the resulting $h^{(1)}(C)$.

Solution

ENG)

- Write these three sentences.
 - In FastGCN the sampling distribution is fixed in advance and does not depend on the parent nodes.
 - In ASGCN the proposal distribution is conditioned on the nodes of the upper layer.
 - This makes almost every sampled node useful and reduces the variance of the estimate.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=3$, $d(D)=2$, $d(E)=3$

Divide each row by its degree. $\tilde{A} = [0, 0, 0.5, 0, 0.5]; [0, 0, 0, 0.5, 0.5]; [0.3333, 0, 0, 0.3333, 0.3333]; [0, 0.5, 0.5, 0, 0]; [0.3333, 0.3333, 0.3333, 0, 0]$ (each row sums to 1)

The parent node is C, so only its row of $\tilde{A}$ matters.

$N(C) = \text{sum of that row} = 1$. With row normalisation it is 1.

$$p(A|C) = \tilde{A}_{CA} / N = 0.3333$$

$$p(D|C) = \tilde{A}_{CD} / N = 0.3333$$

$$p(E|C) = \tilde{A}_{CE} / N = 0.3333$$

$$b) N(C) = 1, p(A|C) = 0.3333, p(D|C) = 0.3333, p(E|C) = 0.3333$$

The expectation is the sum of probability times value: $0.3333 \times -1 + 0.3333 \times 1 + 0.3333 \times -2 = -0.6667$

$N = 1$ and $W = 1$, so the value is unchanged. It equals the plain neighbour mean.

ReLU clips the negatives: $-0.6667 \rightarrow 0$

$$c) E = -0.6667, h_1(C) = 0$$

4. (APPNP) Given a graph with an adjacency matrix A and a prediction matrix H produced by an MLP as follows:

$$A = \begin{bmatrix} 0 & 0 & 1 & 1 \\ 0 & 0 & 1 & 1 \\ 1 & 1 & 0 & 0 \\ 1 & 1 & 0 & 0 \end{bmatrix}$$

$$H = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$$

APPNP propagates the predictions with a teleport probability α :

$$Z^{(0)} = H, \quad Z^{(k+1)} = (1 - \alpha) \tilde{A} Z^{(k)} + \alpha H$$

where $\tilde{A} = D^{-1}A$ and $\alpha = 0.5$. No nonlinearity is applied between steps.

- Explain the teleport mechanism and why it prevents over-smoothing.
- Calculate $\tilde{A}$ and $Z^{(1)}$.
- Calculate $Z^{(2)}$.

Solution

ENG)

a) Write these three sentences.

- At every propagation step the walk returns to the root node with probability alpha. This is the teleport mechanism.
- Because a fixed fraction of the mass is pulled back at every step, the distribution stays local.
- So we can use many propagation steps without over-smoothing, and propagation has no extra parameters.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=2$, $d(D)=2$

Divide each row by its degree. $\tilde{A} = [0, 0, 0.5, 0.5]; [0, 0, 0.5, 0.5]; [0.5, 0.5, 0, 0]; [0.5, 0.5, 0, 0]$ (each row sums to 1)

b) $\tilde{A} = [0, 0, 0.5, 0.5]; [0, 0, 0.5, 0.5]; [0.5, 0.5, 0, 0]; [0.5, 0.5, 0, 0]$

Start from $Z_0 = H = [0, 0, 0, 1]$

Propagate one step: $\tilde{A} Z_0 = [0.5, 0.5, 0, 0]$

Multiply by $(1-\alpha) = 0.5$: $[0.25, 0.25, 0, 0]$

Add $\alpha H = [0, 0, 0, 0.5]$. This is the teleport term.

$$Z_1 = [0.25, 0.25, 0, 0.5]$$

One more step: $\tilde{A} Z_1 = [0.25, 0.25, 0.25, 0.25]$

Multiply by $(1-\alpha)$ and add αH **again**. It is added at every step.

c) $Z_2 = [0.125, 0.125, 0.125, 0.625]$

5. (GPR-GNN) Given a graph with an adjacency matrix A and hidden features $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \\ 1 & 0 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 4 \\ -2 \\ -1 \\ 4 \end{bmatrix}$$

GPR-GNN separates feature extraction from topological propagation:

$$H^{(k)} = \tilde{A} H^{(k-1)}, \quad Z = \sum_{k=0}^K \gamma_k H^{(k)}$$

where $\tilde{A} = D^{-1}A$, $K = 2$, $\gamma_0 = 0.5$, $\gamma_1 = 0.2$, $\gamma_2 = -0.2$.

- What can GPR-GNN do that APPNP cannot, and why do negative γ_k matter?
- Calculate $H^{(1)}$ and $H^{(2)}$.
- Calculate Z .

Solution

ENG)

- Write these three sentences.
 - APPNP fixes the coefficients to alpha times (1-alpha) to the power k, which are always positive and decreasing.
 - GPR-GNN learns the weights, and they are allowed to be negative.
 - A negative weight subtracts that hop, so the model can act as a high-pass filter on heterophilic graphs.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=2$, $d(D)=2$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0, 0.5]; [0.5, 0, 0.5, 0]; [0, 0.5, 0, 0.5]; [0.5, 0, 0.5, 0]$ (each row sums to 1)

H_0 stays as it is: $[4, -2, -1, 4]$

$H_1 = \tilde{A} H_0 = [1, 1.5, 1, 1.5]$

$H_2 = \tilde{A} H_1 = [1.5, 1, 1.5, 1]$

b) $H_1 = [1, 1.5, 1, 1.5]$, $H_2 = [1.5, 1, 1.5, 1]$

Now take the weighted sum with gamma. There is no ReLU during propagation.

$\gamma_0 H_0 = 0.5 \times [4, -2, -1, 4] = [2, -1, -0.5, 2]$

$\gamma_1 H_1 = 0.2 \times [1, 1.5, 1, 1.5] = [0.2, 0.3, 0.2, 0.3]$

$\gamma_2 H_2 = -0.2 \times [1.5, 1, 1.5, 1] = [-0.3, -0.2, -0.3, -0.2]$

γ_2 is negative, so that hop is subtracted. Dropping the sign gives a wrong answer.

Add the three vectors element-wise.

c) $Z = [1.9, -0.9, -0.6, 2.1]$

6. (MixHop) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 2 \\ -2 \\ 1 \\ -1 \end{bmatrix}$$

A MixHop layer concatenates several powers of the adjacency matrix:

$$H^{(i+1)} = \parallel_{j \in P} \sigma(\hat{A}^j H^{(i)} W_j^{(i)})$$

where $P = \{0, 1, 2\}$, $\hat{A} = D^{-1}A$, $W_j^{(i)} = 1$ and $\text{ReLU}(x) = \max(0, x)$.

- State Theorem 1 of MixHop about the two-hop Delta Operator.
- Calculate $\hat{A}^0 H^{(0)}$, $\hat{A}^1 H^{(0)}$ and $\hat{A}^2 H^{(0)}$.
- Calculate $H^{(1)}$ and the Delta Operator $\sigma(\hat{A} H^{(0)}) - \sigma(\hat{A}^2 H^{(0)})$.

Solution

ENG)

a) Write these three sentences.

- Theorem 1 says that a vanilla GCN cannot represent the two-hop Delta Operator, even with infinitely many layers.
- Stacking GCN layers collapses to $\hat{A}$ to the power l times W , so the intermediate hops are never kept separately.
- MixHop computes each power separately and concatenates them, so a later layer can subtract one from another.

Degrees: $d(A)=2$, $d(B)=1$, $d(C)=2$, $d(D)=3$

Divide each row by its degree. $\tilde{A} = [0, 0, 0.5, 0.5]; [0, 0, 0, 1]; [0.5, 0, 0, 0.5]; [0.3333, 0.3333, 0.3333, 0]$ (each row sums to 1)

$\hat{A}^0$ is the identity, so it is just H^0 : $[2, -2, 1, -1]$

$\hat{A}^1 H^0 = [0, -1, 0.5, 0.3333]$ (one hop, the neighbour mean)

$\hat{A}^2 H^0 = \tilde{A}(\tilde{A} H^0) = [0.4167, 0.3333, 0.1667, -0.1667]$ (two hops)

ReLU is applied to each block separately, not after concatenation.

After ReLU: $[2, 0, 1, 0] \parallel [0, 0, 0.5, 0.3333] \parallel [0.4167, 0.3333, 0.1667, 0]$

The delta operator subtracts the two-hop block from the one-hop block.

b) $A^0 H = [2, -2, 1, -1]$, $A^1 H = [0, -1, 0.5, 0.3333]$, $A^2 H = [0.4167, 0.3333, 0.1667, -0.1667]$

c) $H^1 =$ concatenation of the three blocks: $[2, 0, 1, 0] \parallel [0, 0, 0.5, 0.3333] \parallel [0.4167, 0.3333, 0.1667, 0]$

Delta = $[-0.4167, -0.3333, 0.3333, 0.3333]$

7. (**PairNorm**) After one GCN layer the node features of a four-node graph are

$$X = \begin{bmatrix} 4 & 0 \\ 1 & 3 \\ 1 & -1 \\ -1 & 4 \end{bmatrix}$$

PairNorm normalizes them with a centering step and a scaling step:

$$x_i^c = x_i - (1/n) \sum_i x_i \quad (\text{Center})$$

$$x_i = s \cdot x_i^c / \text{sqrt}((1/n) \sum_i \|x_i^c\|_2^2) \quad (\text{Scale})$$

with $s = 1$ and $n = 4$.

- Explain node-wise and feature-wise over-smoothing, and what PairNorm keeps constant.
- Calculate the centered features.
- Calculate the scaling denominator and the final x_i .

Solution

ENG)

- Write these three sentences.
 - Node-wise over-smoothing means the features of connected nodes become too similar.
 - Feature-wise over-smoothing means all node features become indistinguishable across the whole graph.
 - PairNorm keeps the total pairwise squared distance between node representations constant across layers.

First take the mean over the nodes, separately for each feature column.

Column 1: $(4 + 1 + 1 + -1)/4 = 1.25$

Column 2: $(0 + 3 + -1 + 4)/4 = 1.5$

Subtract the mean from every row. Check: each column now sums to zero.

Node 1: squared norm = 9.8125

Node 2: squared norm = 2.3125

Node 3: squared norm = 6.3125

Node 4: squared norm = 11.3125

Mean = $(9.8125 + 2.3125 + 6.3125 + 11.3125)/4 = 7.4375$

Take the square root: $\text{sqrt}(7.4375) = 2.7272$. Divide every row by this single number.

b) Mean = [1.25, 1.5], centered = [2.75, -1.5]; [-0.25, 1.5]; [-0.25, -2.5]; [-2.25, 2.5]

c) Scaling denominator = 2.7272, $x = [1.0084, -0.55]; [-0.0917, 0.55]; [-0.0917, -0.9167]; [-0.825, 0.9167]$

8. (DropEdge) Consider the complete graph K_4 on nodes A, B, C and D with the adjacency matrix A and initial node feature matrix $H^{(0)}$ below.

$$A = \begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 0 \\ -3 \\ 4 \\ 3 \end{bmatrix}$$

DropEdge keeps each undirected edge independently with probability $1 - p$ at every epoch:

$$H^{(1)} = \sigma((D^{(t)})^{-1} A^{(t)} H^{(0)})$$

where $A^{(t)}$ and $D^{(t)}$ are rebuilt from the surviving edges only, $p = 1/3$, and $\text{ReLU}(x) = \max(0, x)$.

- Which two problems does DropEdge target, and how does it differ from Dropout?
- Calculate the expected number of surviving edges, and $H^{(1)}$ on the full graph.
- Suppose the edges (A, C) and (B, D) are dropped. Recalculate $H^{(1)}$.

Solution

ENG)

- Write these three sentences.
 - DropEdge targets over-fitting and over-smoothing.
 - A different random subgraph is used at every epoch, which works as data augmentation on the graph.
 - Dropout zeroes entries of the feature tensor, while DropEdge removes edges of the adjacency matrix.

K_4 has 6 edges. Each survives with probability $1 - 1/3 = 2/3$, so the expectation is 4.

First the full graph: every node has degree 3.

A: neighbours B, C, D. $(-3 + 4 + 3)/3 = 1.3333$

B: neighbours A, C, D. $(0 + 4 + 3)/3 = 2.3333$

C: neighbours A, B, D. $(0 + -3 + 3)/3 = 0$

D: neighbours A, B, C. $(0 + -3 + 4)/3 = 0.3333$

Now drop the two given edges and **recount the degrees**. This is where most mistakes happen.

Degrees after dropping: $d(A)=2$, $d(B)=2$, $d(C)=2$, $d(D)=2$

A: neighbours B, D. $(-3 + 3)/2 = 0$

B: neighbours A, C. $(0 + 4)/2 = 2$

C: neighbours B, D. $(-3 + 3)/2 = 0$

D: neighbours A, C. $(0 + 4)/2 = 2$

b) Expected number of surviving edges = 4, $H_1 = [1.3333, 2.3333, 0, 0.3333]$

c) $H_1 = [0, 2, 0, 2]$

9. (Node2vec) Consider an undirected graph G given in the following figure. A random walk has just moved from node A to node B , with $p = 2$ and $q = 0.5$.




Figure 1. Graph G

The unnormalized transition probability from B to a next node x is $\pi_{Bx} = \alpha_{pq}(A, x)$, where $\alpha = 1/p$ if $d(A, x) = 0$, $\alpha = 1$ if $d(A, x) = 1$, and $\alpha = 1/q$ if $d(A, x) = 2$.

- Explain BFS and DFS, and explain which one a small q makes the walk behave like.
- List all possible next nodes from node B and calculate π_{Bx} .
- Calculate the normalized transition probabilities.

Solution

ENG)

- Write these three sentences.
 - BFS explores the immediate neighbourhood first, so the embedding captures the structural role of a node.
 - DFS moves far away from the source, so the embedding captures community membership.
 - A small q makes $1/q$ large, so the walk moves outward and behaves like DFS.

The previous node is A and the walk is now at B . All distances are measured from A .

A : $x = A$ is the previous node, $d = 0$, $\alpha = 1/p = 0.5$

C : x is two hops from A , $d = 2$, $\alpha = 1/q = 2$

D : x is also a neighbour of A , $d = 1$, $\alpha = 1 = 1$

Sum them: $0.5 + 2 + 1 = 3.5$

b) $\pi(B \rightarrow A) = 0.5$, $\pi(B \rightarrow C) = 2$, $\pi(B \rightarrow D) = 1$

c) $P(A) = 0.1429$, $P(C) = 0.5714$, $P(D) = 0.2857$

10. (LINE) Consider an undirected graph G with node embeddings

$$u_A = 1, u_B = 0.5, u_C = -0.5, u_D = -1$$

and context embeddings

$$u'_A = -0.5, u'_B = 0.5, u'_C = 1, u'_D = -1.$$

The second-order proximity is $p_2(j|x) = \exp(u'_j u_x) / \sum_k \exp(u'_k u_x)$. Given $e^{0.25} = 1.28$, $e^{0.5} = 1.65$, $e^{-0.25} = 0.78$, $e^{-0.5} = 0.61$.

- What is second-order proximity, and how does it differ from first-order proximity?
- Calculate $p_2(A|B)$, $p_2(B|B)$, $p_2(C|B)$ and $p_2(D|B)$.

Solution

ENG)

- Write these three sentences.

- First-order proximity is the similarity between two nodes that are directly connected by an edge.
- Second-order proximity is the similarity between two nodes that share many neighbours, even without an edge.
- So LINE gives each node two vectors: its own embedding and its context embedding.

The base node is B with $u(B) = 0.5$. The other side uses the context vector u' .

$$A: u'(A) u(B) = -0.5 \times 0.5 = -0.25, \exp = 0.78$$

$$B: u'(B) u(B) = 0.5 \times 0.5 = 0.25, \exp = 1.28$$

$$C: u'(C) u(B) = 1 \times 0.5 = 0.5, \exp = 1.65$$

$$D: u'(D) u(B) = -1 \times 0.5 = -0.5, \exp = 0.61$$

The denominator is the sum of all four: $0.78 + 1.28 + 1.65 + 0.61 = 4.32$

$$b) p_2(A|B) = 0.1806, p_2(B|B) = 0.2963, p_2(C|B) = 0.3819, p_2(D|B) = 0.1412$$

11. (GCN) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 1 & 1 \\ 1 & 1 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 1 \\ -2 \\ -1 \\ 4 \\ 4 \end{bmatrix}$$

Assume that the hidden layer of a GCN model of all nodes at layer (k) can be calculated as:

$$H^{(k)} = \sigma(\tilde{A} \cdot H^{(k-1)}),$$

where $\tilde{A} = D^{-1}A$ and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- Explain how information is propagated in a GCN layer.
- Calculate $\tilde{A}$.
- Calculate the output of the GCN model at layer $k = 1$.

Solution

ENG)

a) Write these three sentences.

- In one layer each node aggregates the features of its direct neighbours and takes the average.
- The row-normalized matrix $\tilde{A}$ equals $D^{-1}A$, which performs this averaging.
- After k layers a node has received information from all nodes within k hops.

Degrees: $d(A)=2$, $d(B)=3$, $d(C)=3$, $d(D)=1$, $d(E)=3$

Divide each row by its degree. $\tilde{A} = [0, 0, 0.5, 0, 0.5]; [0, 0, 0.3333, 0.3333, 0.3333]; [0.3333, 0.3333, 0, 0, 0.3333]; [0, 1, 0, 0, 0]; [0.3333, 0.3333, 0.3333, 0, 0]$ (each row sums to 1)

b) $\tilde{A} = [0, 0, 0.5, 0, 0.5]; [0, 0, 0.3333, 0.3333, 0.3333]; [0.3333, 0.3333, 0, 0, 0.3333]; [0, 1, 0, 0, 0]; [0.3333, 0.3333, 0.3333, 0, 0]$

Multiply $\tilde{A}$ by H^0 . Each entry is the mean of that node's neighbours.

A: neighbours C, E. $(-1 + 4)/2 = 1.5$

B: neighbours C, D, E. $(-1 + 4 + 4)/3 = 2.3333$

C: neighbours A, B, E. $(1 + -2 + 4)/3 = 1$

D: neighbours B. $(-2)/1 = -2$

E: neighbours A, B, C. $(1 + -2 + -1)/3 = -0.6667$

ReLU clips the negatives: $-2 \rightarrow 0$, $-0.6667 \rightarrow 0$

c) $H_1 = [1.5, 2.3333, 1, 0, 0]$

12. (GraphSAGE) Consider an undirected graph G of five nodes given in the following figure. Each node has initial features that are the numbers standing next to it. According to the GraphSAGE model with an AGGREGATE that is a MEAN function, the feature of a node i at layer k can be updated as:

$$h_{N(i)}^{(k)} = \text{MEAN}(\{ h_u^{(k-1)}, \forall u \in N(i) \})$$

$$h_i^{(k)} = \text{ReLU}(h_i^{(k-1)} \parallel h_{N(i)}^{(k)})$$

where $\parallel$ is a concatenation and $\text{ReLU}(x) = \max(0, x)$.




Figure. Graph G

- What is the main idea behind GraphSAGE, and how does it differ from traditional GCNs?
- Calculate the feature of each node at $k = 1$.
- Calculate a graph-level embedding h_G by using a 'Mean' global pooling when $k = 1$.

Solution

ENG)

a) Write these three sentences.

- GraphSAGE learns an aggregator function instead of a fixed embedding for each node.
- It samples a fixed number of neighbours, so it can be trained in mini-batches on very large graphs.
- A GCN is transductive, but GraphSAGE is inductive, so it can embed a node that was never seen in training.

First compute the neighbour mean. The node itself is not included.

A: neighbours B, D, E. $(3 + -2 + 2)/3 = 1$

B: neighbours A, C, D. $(0 + 4 + -2)/3 = 0.6667$

C: neighbours B, D. $(3 + -2)/2 = 0.5$

D: neighbours A, B, C. $(0 + 3 + 4)/3 = 2.3333$

E: neighbours A. $(0)/1 = 0$

Concatenate the node's own value with the neighbour mean. This is not an addition.

ReLU clips the negatives: $-2 \rightarrow 0$

The graph embedding is the mean per position: first $7 / 5$, second $4.5 / 5$

b) $h(A) = [0, 1]$, $h(B) = [3, 0.6667]$, $h(C) = [4, 0.5]$, $h(D) = [0, 2.3333]$, $h(E) = [2, 0]$

c) $h_G = [1.8, 0.9]$

13. (ClusterGCN) Consider an undirected graph G of six nodes given in the following figure. The graph G contains two clusters G_1 and G_2 . Each node has initial features that are the numbers standing next to it. According to the ClusterGCN model, the feature of a node i at layer k can be updated as:

$$h_{N(i)}^{(k)} = \text{MEAN}(\{ h_u^{(k-1)}, \forall u \in N(i), G_u = G_i \})$$

$$h_i^{(k)} = \text{ReLU}(h_i^{(k-1)} \parallel h_{N(i)}^{(k)})$$

where $\parallel$ is a concatenation and $\text{ReLU}(x) = \max(0, x)$.




Figure. Graph G

- What is mini-batch training, and why does ClusterGCN need it?
- Calculate the output representations of all nodes at layer $k = 1$.

Solution

ENG)

- Write these three sentences.

- Mini-batch training updates the parameters using a small subset of the data at a time.
- A large graph does not fit in memory, so ClusterGCN partitions the graph into densely connected clusters.
- A good partition keeps most edges inside a batch, so little information is lost at the boundary.

Edges that cross a group are treated as absent. Cross them out first.

A: no neighbour inside the group, so the mean is 0.

B: neighbours C. $(1)/1 = 1$

C: neighbours B. $(3)/1 = 3$

D: neighbours E. $(4)/1 = 4$

E: neighbours D, F. $(2 + 2)/2 = 2$

F: neighbours E. $(4)/1 = 4$

Concatenate the neighbour mean after the node's own value, then apply ReLU to both entries.

- $h(A) = [0, 0]$, $h(B) = [3, 1]$, $h(C) = [1, 3]$, $h(D) = [2, 4]$, $h(E) = [4, 2]$, $h(F) = [2, 4]$

14. (GraphSAINT) Consider an undirected graph G of six nodes given in the following figure. The graph G has a sampled subgraph G_1 . Each node has initial features that are the numbers standing next to it. According to the GraphSAINT model, the feature of a node i at layer k can be updated as:

$$h_{N(i)}^{(k)} = \text{MEAN}(\{ h_u^{(k-1)}, \forall u \in N(i), G_u = G_i \})$$

$$h_i^{(k)} = \text{ReLU}(h_i^{(k-1)} \parallel h_{N(i)}^{(k)})$$

where $\parallel$ is a concatenation and $\text{ReLU}(x) = \max(0, x)$.




Figure. Graph G

- How does GraphSAINT's sampling differ from node-wise or layer-wise sampling?
- Calculate the output representations of nodes A, B, C and D at layer $k = 1$.

Solution

ENG)

- Write these three sentences.
 - GraphSAGE and FastGCN sample while building the layers, so sampling is repeated at every layer.
 - GraphSAINT samples a subgraph first and then builds the full GNN on that subgraph.
 - So the cost does not grow with depth, and normalization coefficients correct the sampling bias.

Edges that leave the sampled subgraph are discarded.

A: neighbours C. $(-1)/1 = -1$

B: neighbours C. $(-1)/1 = -1$

C: neighbours A, B. $(1 + -1)/2 = 0$

D: no neighbour inside the group, so the mean is 0.

Concatenate the neighbour mean after the node's own value, then apply ReLU.

b) $h(A) = [1, 0]$, $h(B) = [0, 0]$, $h(C) = [0, 0]$, $h(D) = [2, 0]$

15. (JK Network) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 1 & 1 \\ 1 & 0 & 0 & 1 & 1 \\ 1 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 4 \\ -3 \\ -3 \\ -3 \\ 3 \end{bmatrix}$$

Assume that the output of a JK network model of all nodes at layer (k) can be calculated as:

$$H^{(k)} = \max(\sigma(\tilde{A} H^{(0)}), \sigma(\tilde{A} H^{(1)}), \dots, \sigma(\tilde{A} H^{(k-1)}))$$

where $\tilde{A} = D^{-1}A$ and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- Explain how Jumping Knowledge Networks help mitigate the over-smoothing problem.
- Calculate $\tilde{A}$ and $H^{(1)}$.
- Calculate the output representations at layer $k = 2$.

Solution

ENG)

a) Write these three sentences.

- In a deep GCN every node repeatedly averages its neighbours, so all representations become similar. This is over-smoothing.
- JKNet keeps the representation of every layer and combines them at the end by max, concatenation or an LSTM.
- The best depth is different for each node, so each node can select the depth that suits it.

Degrees: $d(A)=4$, $d(B)=3$, $d(C)=1$, $d(D)=2$, $d(E)=2$

Divide each row by its degree. $\text{Atilde} = [0, 0.25, 0.25, 0.25, 0.25]; [0.3333, 0, 0, 0.3333, 0.3333]; [1, 0, 0, 0, 0]; [0.5, 0.5, 0, 0, 0]; [0.5, 0.5, 0, 0, 0]$ (each row sums to 1)

b) $\text{Atilde} = [0, 0.25, 0.25, 0.25, 0.25]; [0.3333, 0, 0, 0.3333, 0.3333]; [1, 0, 0, 0, 0]; [0.5, 0.5, 0, 0, 0]; [0.5, 0.5, 0, 0, 0]$

First build $H1 = \text{ReLU}(\text{Atilde } H0)$

A: neighbours B, C, D, E. $(-3 + -3 + -3 + 3)/4 = -1.5$

B: neighbours A, D, E. $(4 + -3 + 3)/3 = 1.3333$

C: neighbours A. $(4)/1 = 4$

D: neighbours A, B. $(4 + -3)/2 = 0.5$

E: neighbours A, B. $(4 + -3)/2 = 0.5$

ReLU clips the negatives: $-1.5 \rightarrow 0$

$H1 = [0, 1.3333, 4, 0.5, 0.5]$

For $k = 2$ there are two terms. First: $\text{ReLU}(\text{Atilde } H0) = [0, 1.3333, 4, 0.5, 0.5]$

Second: $\text{ReLU}(\text{Atilde } H1) = [1.5833, 0.3333, 0, 0.6667, 0.6667]$

The max is taken element-wise, not by choosing one whole vector.

c) $H2 = [1.5833, 1.3333, 4, 0.6667, 0.6667]$

16. (GCNII) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 0 & 0 \\ 1 & 0 & 0 & 1 & 1 \\ 1 & 0 & 0 & 1 & 1 \\ 0 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} -3 \\ 3 \\ 3 \\ 0 \\ -1 \end{bmatrix}$$

Assume that the output of a GCNII model of all nodes at layer (k) can be calculated as:

$$H^{(k)} = \sigma[((1 - \beta) I_n) \cdot ((1 - \alpha) \tilde{A} \cdot H^{(k-1)} + \alpha H^{(0)})]$$

where $\tilde{A} = D^{-1}A$, I_n is the identity matrix, $\alpha = 0.2$, $\beta = 0.4$, and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- What are the two major components introduced in GCNII to overcome the depth limitation?
- Calculate $\tilde{A}$.
- Calculate the output representations at layer $k = 1$.

Solution

ENG)

- Write these three sentences.
 - The two components are the initial residual connection and the identity mapping.
 - The initial residual connection adds a fraction alpha of the input H zero back at every layer.
 - The identity mapping keeps the weight matrix close to the identity, so a deep model is not destroyed.

Degrees: $d(A)=2$, $d(B)=3$, $d(C)=3$, $d(D)=2$, $d(E)=2$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0.5, 0, 0]; [0.3333, 0, 0, 0.3333, 0.3333]; [0.3333, 0, 0, 0.3333, 0.3333]; [0, 0.5, 0.5, 0, 0]; [0, 0.5, 0.5, 0, 0]$ (each row sums to 1)

b) $\tilde{A} = [0, 0.5, 0.5, 0, 0]; [0.3333, 0, 0, 0.3333, 0.3333]; [0.3333, 0, 0, 0.3333, 0.3333]; [0, 0.5, 0.5, 0, 0]; [0, 0.5, 0.5, 0, 0]$

Start with the inner bracket. $\tilde{A} H^0 = [3, -1.3333, -1.3333, 3, 3]$

Multiply by $(1-\alpha) = 0.8$: $[2.4, -1.0667, -1.0667, 2.4, 2.4]$

Add $\alpha H^0 = 0.2 \times [-3, 3, 3, 0, -1] = [-0.6, 0.6, 0.6, 0, -0.2]$

Inner bracket = $[1.8, -0.4667, -0.4667, 2.4, 2.2]$

Multiply by the outer $(1-\beta) = 0.6$. Dropping this factor scales the answer.

ReLU clips the negatives: $-0.28 \rightarrow 0$, $-0.28 \rightarrow 0$

c) $H^1 = [1.08, 0, 0, 1.44, 1.32]$

17. (DeepGCN) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} -3 \\ 1 \\ 2 \\ 0 \end{bmatrix}$$

Assume that the output of a DeepGCN model with a residual connection at layer (k) is:

$$H^{(k)} = \sigma(\tilde{A} \cdot H^{(k-1)}) + H^{(k-1)}$$

where $\tilde{A} = D^{-1}A$ and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- Explain why a residual connection allows a GCN to be made deeper.
- Calculate the output representations at layer $k = 1$ and $k = 2$.
- Compare $H^{(1)}$ with the output of a plain GCN without the residual connection.

Solution

ENG)

- Write these three sentences.
 - Deep GCNs suffer from vanishing gradients and over-smoothing, so adding layers makes the model worse.
 - A residual connection passes the input forward unchanged, so each layer only has to learn the change.
 - If no change is useful the layer can output zero, and gradients also flow back through the skip path.

Degrees: $d(A)=2$, $d(B)=3$, $d(C)=3$, $d(D)=2$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0.5, 0]; [0.3333, 0, 0.3333, 0.3333]; [0.3333, 0.3333, 0, 0.3333]; [0, 0.5, 0.5, 0]$ (each row sums to 1)

$\tilde{A} H^0 = [1.5, -0.3333, -0.6667, 1.5]$

Apply ReLU first: $[1.5, 0, 0, 1.5]$

Then add H^0 . The addition sits outside the ReLU, so negatives can survive.

For $k = 2$, feed H^1 back in and add H^1 , not H^0 .

$\tilde{A} H^1 = [1.5, 0.6667, 0.3333, 1.5]$, after ReLU $[1.5, 0.6667, 0.3333, 1.5]$

b) $H^1 = [-1.5, 1, 2, 1.5]$

$H^2 = [0, 1.6667, 2.3333, 3]$

c) Plain GCN without residual: $H^1 = [1.5, 0, 0, 1.5]$ (the node's own value and negatives disappear)

4회차

1. (Relational Classification) Consider an undirected graph G of five nodes given in the following figure. Nodes 1 and 5 are labelled ($P(Y_1) = 1$, $P(Y_5) = 0$) and the others are unlabelled and initialized to 0.5. The relational classifier updates every unlabelled node as:

$$P(Y_v = c) = (\sum_{(v,u) \in E} A_{v,u} P(Y_u = c)) / (\sum_{(v,u) \in E} A_{v,u})$$




Figure. Graph G

- Explain the key assumption behind relational classification, and state one limitation.
- Calculate $P(Y_v)$ for every unlabelled node after the 1st iteration.
- Calculate $P(Y_v)$ after the 2nd iteration, using the previous iteration's values only.

Solution

ENG)

- Write these three sentences.
 - Relational classification assumes homophily: connected nodes tend to share the same label.
 - It repeatedly sets each unlabelled node to the weighted average of its neighbours' class probabilities.
 - Its limitation is that it uses only the graph structure and ignores node features.

Write the initial values: node 1 is 1, node 5 is 0, the rest start at 0.5.

First iteration. Only the unlabelled nodes are replaced by the mean of their neighbours.

Node 1 is labelled, so it never changes. $P = 1$

Node 2: neighbours 3, 5. $(0.5 + 0)/2 = 0.25$

Node 3: neighbours 2, 4, 5. $(0.5 + 0.5 + 0)/3 = 0.3333$

Node 4: neighbours 1, 3. $(1 + 0.5)/2 = 0.75$

Node 5 is labelled, so it never changes. $P = 0$

b) $P(Y_1) = 1$, $P(Y_2) = 0.25$, $P(Y_3) = 0.3333$, $P(Y_4) = 0.75$, $P(Y_5) = 0$

Second iteration. Use only the values from the first iteration, not the ones just computed.

Node 1 is labelled, so it never changes. $P = 1$

Node 2: neighbours 3, 5. $(0.3333 + 0)/2 = 0.1667$

Node 3: neighbours 2, 4, 5. $(0.25 + 0.75 + 0)/3 = 0.3333$

Node 4: neighbours 1, 3. $(1 + 0.3333)/2 = 0.6667$

Node 5 is labelled, so it never changes. $P = 0$

c) $P(Y_1) = 1$, $P(Y_2) = 0.1667$, $P(Y_3) = 0.3333$, $P(Y_4) = 0.6667$, $P(Y_5) = 0$

2. (FastGCN) Consider the undirected graph G with adjacency matrix A and node feature vector $h^{(0)}$ below. FastGCN replaces the full neighbourhood sum of a GCN layer by a layer-wise Monte-Carlo estimate:

$$A = \begin{bmatrix} 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 1 & 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 & 0 \end{bmatrix}$$

$$h^{(0)} = \begin{bmatrix} 3 \\ 3 \\ 2 \\ 0 \\ 2 \end{bmatrix}$$

$$h^{(l+1)}(v) = \sigma \left(\frac{1}{t_l} \sum_{j=1}^{t_l} \tilde{A}(v, u_j) h^{(l)}(u_j) \right)$$

where $\tilde{A} = D^{-1}A$ and the t_l nodes are drawn once per layer from $q(u) \propto \|\tilde{A}(:, u)\|^2$, with $\text{ReLU}(x) = \max(0, x)$.

- How does layer-wise sampling differ from the node-wise sampling of GraphSAGE?
- Calculate the importance distribution $q(u)$ over the five nodes.
- Suppose the sampled set is $\{D, E\}$ ($t_l = 2$). Calculate the estimate of $h^{(1)}(A)$.

Solution

ENG)

- Write these three sentences.
 - In node-wise sampling, each node samples its own neighbours, so the cost grows exponentially with depth.
 - FastGCN samples a fixed-size set of nodes once per layer, shared by all nodes in that layer.
 - Therefore the cost grows only linearly with depth instead of exponentially.

Degrees: $d(A)=3$, $d(B)=2$, $d(C)=2$, $d(D)=3$, $d(E)=2$

Divide each row by its degree. $\tilde{A} = [0, 0, 0.3333, 0.3333, 0.3333]; [0, 0, 0.5, 0.5, 0]; [0.5, 0.5, 0, 0, 0]; [0.3333, 0.3333, 0, 0, 0.3333]; [0.5, 0, 0, 0.5, 0]$ (each row sums to 1)

$q(u)$ uses the squared norm of column u of $\tilde{A}$, not of a row.

Column A: squared sum = 0.6111

Column B: squared sum = 0.3611

Column C: squared sum = 0.3611

Column D: squared sum = 0.6111

Column E: squared sum = 0.2222

Divide by the total 2.1667 to turn them into probabilities.

b) $q(A) = 0.2821$, $q(B) = 0.1667$, $q(C) = 0.1667$, $q(D) = 0.2821$, $q(E) = 0.1026$

For (c) look only at the sampled nodes: read $\tilde{A}(A, u)$ for each.

$\tilde{A}(A, D) h(D) = 0.3333 \times 0 = 0$

$\tilde{A}(A, E) h(E) = 0.3333 \times 2 = 0.6667$

Add them and divide by the number sampled (2), not by the degree: $(0 + 0.6667)/2 = 0.3333$

No negative value, so ReLU changes nothing.

c) $h^1(A) = 0.3333$

3. (ASGCN) Consider the undirected graph G with adjacency matrix A and node feature vector $h^{(0)}$ below. ASGCN rewrites a GCN layer as an expectation:

$$A = \begin{bmatrix} 0 & 1 & 1 & 0 & 0 \\ 1 & 0 & 1 & 1 & 0 \\ 1 & 1 & 0 & 1 & 1 \\ 0 & 1 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \end{bmatrix}$$

$$h^{(0)} = \begin{bmatrix} 0 \\ -1 \\ 2 \\ 4 \\ 2 \end{bmatrix}$$

$$h^{(l+1)}(v_i) = \sigma(W^{(l)} N(v_i) E_{p(u_j | v_i)}[h^{(l)}(u_j)])$$

where $\tilde{A} = D^{-1}A$, $N(v_i) = \sum_j \tilde{A}_{ij}$, $p(u_j | v_i) = \tilde{A}_{ij} / N(v_i)$, $W^{(l)} = 1$ and $\text{ReLU}(x) = \max(0, x)$.

- What does the proposal distribution depend on in ASGCN that it does not in FastGCN?
- Calculate $N(v_i)$ and $p(u_j | v_i)$ for the parent node $v_i = B$.
- Calculate $E_{p(u_j | B)}[h^{(0)}(u_j)]$ and the resulting $h^{(1)}(B)$.

Solution

ENG)

- Write these three sentences.

- In FastGCN the sampling distribution is fixed in advance and does not depend on the parent nodes.
- In ASGCN the proposal distribution is conditioned on the nodes of the upper layer.
- This makes almost every sampled node useful and reduces the variance of the estimate.

Degrees: $d(A)=2$, $d(B)=3$, $d(C)=4$, $d(D)=2$, $d(E)=1$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0.5, 0, 0]; [0.3333, 0, 0.3333, 0.3333, 0]; [0.25, 0.25, 0, 0.25, 0.25]; [0, 0.5, 0.5, 0, 0]; [0, 0, 1, 0, 0]$ (each row sums to 1)

The parent node is B, so only its row of $\tilde{A}$ matters.

$N(B) = \text{sum of that row} = 1$. With row normalisation it is 1.

$$p(A|B) = \tilde{A}_{B,A} / N = 0.3333$$

$$p(C|B) = \tilde{A}_{B,C} / N = 0.3333$$

$$p(D|B) = \tilde{A}_{B,D} / N = 0.3333$$

$$\text{b) } N(B) = 1, p(A|B) = 0.3333, p(C|B) = 0.3333, p(D|B) = 0.3333$$

The expectation is the sum of probability times value: $0.3333 \times 0 + 0.3333 \times 2 + 0.3333 \times 4 = 2$

$N = 1$ and $W = 1$, so the value is unchanged. It equals the plain neighbour mean.

No negative value, so ReLU changes nothing.

$$\text{c) } E = 2, h^{(1)}(B) = 2$$

4. (APPNP) Given a graph with an adjacency matrix A and a prediction matrix H produced by an MLP as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \end{bmatrix}$$

$$H = \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

APPNP propagates the predictions with a teleport probability α :

$$Z^{(0)} = H, \quad Z^{(k+1)} = (1 - \alpha) \tilde{A} Z^{(k)} + \alpha H$$

where $\tilde{A} = D^{-1}A$ and $\alpha = 0.2$. No nonlinearity is applied between steps.

- Explain the teleport mechanism and why it prevents over-smoothing.
- Calculate $\tilde{A}$ and $Z^{(1)}$.
- Calculate $Z^{(2)}$.

Solution

ENG)

a) Write these three sentences.

- At every propagation step the walk returns to the root node with probability alpha. This is the teleport mechanism.
- Because a fixed fraction of the mass is pulled back at every step, the distribution stays local.
- So we can use many propagation steps without over-smoothing, and propagation has no extra parameters.

Degrees: $d(A)=3$, $d(B)=1$, $d(C)=2$, $d(D)=2$

Divide each row by its degree. $\tilde{A} = [0, 0.3333, 0.3333, 0.3333]; [1, 0, 0, 0]; [0.5, 0, 0, 0.5]; [0.5, 0, 0.5, 0]$ (each row sums to 1)

b) $\tilde{A} = [0, 0.3333, 0.3333, 0.3333]; [1, 0, 0, 0]; [0.5, 0, 0, 0.5]; [0.5, 0, 0.5, 0]$

Start from $Z_0 = H = [1, 0, 0, 0]$

Propagate one step: $\tilde{A} Z_0 = [0, 1, 0.5, 0.5]$

Multiply by $(1-\alpha) = 0.8$: $[0, 0.8, 0.4, 0.4]$

Add $\alpha H = [0.2, 0, 0, 0]$. This is the teleport term.

$Z_1 = [0.2, 0.8, 0.4, 0.4]$

One more step: $\tilde{A} Z_1 = [0.5333, 0.2, 0.3, 0.3]$

Multiply by $(1-\alpha)$ and add αH **again**. It is added at every step.

c) $Z_2 = [0.6267, 0.16, 0.24, 0.24]$

5. (GPR-GNN) Given a graph with an adjacency matrix A and hidden features $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 0 \\ 0 & 1 & 0 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} -1 \\ 3 \\ -2 \\ 3 \end{bmatrix}$$

GPR-GNN separates feature extraction from topological propagation:

$$H^{(k)} = \tilde{A} H^{(k-1)}, \quad Z = \sum_{k=0}^K \gamma_k H^{(k)}$$

where $\tilde{A} = D^{-1}A$, $K = 2$, $\gamma_0 = 0.5$, $\gamma_1 = 0.4$, $\gamma_2 = -0.2$.

- What can GPR-GNN do that APPNP cannot, and why do negative γ_k matter?
- Calculate $H^{(1)}$ and $H^{(2)}$.
- Calculate Z .

Solution

ENG)

- Write these three sentences.
 - APPNP fixes the coefficients to alpha times (1-alpha) to the power k, which are always positive and decreasing.
 - GPR-GNN learns the weights, and they are allowed to be negative.
 - A negative weight subtracts that hop, so the model can act as a high-pass filter on heterophilic graphs.

Degrees: $d(A)=2$, $d(B)=3$, $d(C)=2$, $d(D)=1$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0.5, 0]; [0.3333, 0, 0.3333, 0.3333]; [0.5, 0.5, 0, 0]; [0, 1, 0, 0]$
(each row sums to 1)

H_0 stays as it is: $[-1, 3, -2, 3]$

$H_1 = \tilde{A} H_0 = [0.5, 0, 1, 3]$

$H_2 = \tilde{A} H_1 = [0.5, 1.5, 0.25, 0]$

b) $H_1 = [0.5, 0, 1, 3]$, $H_2 = [0.5, 1.5, 0.25, 0]$

Now take the weighted sum with gamma. There is no ReLU during propagation.

$\gamma_0 H_0 = 0.5 \times [-1, 3, -2, 3] = [-0.5, 1.5, -1, 1.5]$

$\gamma_1 H_1 = 0.4 \times [0.5, 0, 1, 3] = [0.2, 0, 0.4, 1.2]$

$\gamma_2 H_2 = -0.2 \times [0.5, 1.5, 0.25, 0] = [-0.1, -0.3, -0.05, 0]$

γ_2 is negative, so that hop is subtracted. Dropping the sign gives a wrong answer.

Add the three vectors element-wise.

c) $Z = [-0.4, 1.2, -0.65, 2.7]$

6. (MixHop) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} -2 \\ -2 \\ -2 \\ 3 \end{bmatrix}$$

A MixHop layer concatenates several powers of the adjacency matrix:

$$H^{(i+1)} = \parallel_{j \in P} \sigma(\hat{A}^j H^{(i)} W_j^{(i)})$$

where $P = \{0, 1, 2\}$, $\hat{A} = D^{-1}A$, $W_j^{(i)} = 1$ and $\text{ReLU}(x) = \max(0, x)$.

- State Theorem 1 of MixHop about the two-hop Delta Operator.
- Calculate $\hat{A}^0 H^{(0)}$, $\hat{A}^1 H^{(0)}$ and $\hat{A}^2 H^{(0)}$.
- Calculate $H^{(1)}$ and the Delta Operator $\sigma(\hat{A} H^{(0)}) - \sigma(\hat{A}^2 H^{(0)})$.

Solution

ENG)

a) Write these three sentences.

- Theorem 1 says that a vanilla GCN cannot represent the two-hop Delta Operator, even with infinitely many layers.
- Stacking GCN layers collapses to $\hat{A}$ to the power 1 times W , so the intermediate hops are never kept separately.
- MixHop computes each power separately and concatenates them, so a later layer can subtract one from another.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=2$, $d(D)=2$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0.5, 0]; [0.5, 0, 0, 0.5]; [0.5, 0, 0, 0.5]; [0, 0.5, 0.5, 0]$ (each row sums to 1)

$\hat{A}^0$ is the identity, so it is just $H^0: [-2, -2, -2, 3]$

$\hat{A}^1 H^0 = [-2, 0.5, 0.5, -2]$ (one hop, the neighbour mean)

$\hat{A}^2 H^0 = \tilde{A}(\tilde{A} H^0) = [0.5, -2, -2, 0.5]$ (two hops)

ReLU is applied to each block separately, not after concatenation.

After ReLU: $[0, 0, 0, 3] \parallel [0, 0.5, 0.5, 0] \parallel [0.5, 0, 0, 0.5]$

The delta operator subtracts the two-hop block from the one-hop block.

b) $A^0 H = [-2, -2, -2, 3]$, $A^1 H = [-2, 0.5, 0.5, -2]$, $A^2 H = [0.5, -2, -2, 0.5]$

c) $H^1 =$ concatenation of the three blocks: $[0, 0, 0, 3] \parallel [0, 0.5, 0.5, 0] \parallel [0.5, 0, 0, 0.5]$

Delta = $[-0.5, 0.5, 0.5, -0.5]$

7. (**PairNorm**) After one GCN layer the node features of a four-node graph are

$$X = \begin{bmatrix} -1 & 0 \\ 0 & 0 \\ 4 & 0 \\ -2 & 0 \end{bmatrix}$$

PairNorm normalizes them with a centering step and a scaling step:

$$x_i^c = x_i - (1/n) \sum_i x_i \quad (\text{Center})$$

$$x_i = s \cdot x_i^c / \text{sqrt}((1/n) \sum_i \|x_i^c\|_2^2) \quad (\text{Scale})$$

with $s = 1$ and $n = 4$.

- Explain node-wise and feature-wise over-smoothing, and what PairNorm keeps constant.
- Calculate the centered features.
- Calculate the scaling denominator and the final x_i .

Solution

ENG)

- Write these three sentences.
 - Node-wise over-smoothing means the features of connected nodes become too similar.
 - Feature-wise over-smoothing means all node features become indistinguishable across the whole graph.
 - PairNorm keeps the total pairwise squared distance between node representations constant across layers.

First take the mean over the nodes, separately for each feature column.

Column 1: $(-1 + 0 + 4 + -2)/4 = 0.25$

Column 2: $(0 + 0 + 0 + 0)/4 = 0$

Subtract the mean from every row. Check: each column now sums to zero.

Node 1: squared norm = 1.5625

Node 2: squared norm = 0.0625

Node 3: squared norm = 14.0625

Node 4: squared norm = 5.0625

Mean = $(1.5625 + 0.0625 + 14.0625 + 5.0625)/4 = 5.1875$

Take the square root: $\text{sqrt}(5.1875) = 2.2776$. Divide every row by this single number.

b) Mean = $[0.25, 0]$, centered = $[-1.25, 0]$; $[-0.25, 0]$; $[3.75, 0]$; $[-2.25, 0]$

c) Scaling denominator = 2.2776, $x = [-0.5488, 0]$; $[-0.1098, 0]$; $[1.6465, 0]$; $[-0.9879, 0]$

8. (DropEdge) Consider the complete graph K_4 on nodes A, B, C and D with the adjacency matrix A and initial node feature matrix $H^{(0)}$ below.

$$A = \begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} -1 \\ 0 \\ 3 \\ 3 \end{bmatrix}$$

DropEdge keeps each undirected edge independently with probability $1 - p$ at every epoch:

$$H^{(1)} = \sigma((D^{(t)})^{-1} A^{(t)} H^{(0)})$$

where $A^{(t)}$ and $D^{(t)}$ are rebuilt from the surviving edges only, $p = 1/3$, and $\text{ReLU}(x) = \max(0, x)$.

- Which two problems does DropEdge target, and how does it differ from Dropout?
- Calculate the expected number of surviving edges, and $H^{(1)}$ on the full graph.
- Suppose the edges (A, D) and (A, C) are dropped. Recalculate $H^{(1)}$.

Solution

ENG)

- Write these three sentences.
 - DropEdge targets over-fitting and over-smoothing.
 - A different random subgraph is used at every epoch, which works as data augmentation on the graph.
 - Dropout zeroes entries of the feature tensor, while DropEdge removes edges of the adjacency matrix.

K_4 has 6 edges. Each survives with probability $1 - 1/3 = 2/3$, so the expectation is 4.

First the full graph: every node has degree 3.

A: neighbours B, C, D. $(0 + 3 + 3)/3 = 2$

B: neighbours A, C, D. $(-1 + 3 + 3)/3 = 1.6667$

C: neighbours A, B, D. $(-1 + 0 + 3)/3 = 0.6667$

D: neighbours A, B, C. $(-1 + 0 + 3)/3 = 0.6667$

Now drop the two given edges and **recount the degrees**. This is where most mistakes happen.

Degrees after dropping: $d(A)=1$, $d(B)=3$, $d(C)=2$, $d(D)=2$

A: neighbours B. $(0)/1 = 0$

B: neighbours A, C, D. $(-1 + 3 + 3)/3 = 1.6667$

C: neighbours B, D. $(0 + 3)/2 = 1.5$

D: neighbours B, C. $(0 + 3)/2 = 1.5$

b) Expected number of surviving edges = 4, $H_1 = [2, 1.6667, 0.6667, 0.6667]$

c) $H_1 = [0, 1.6667, 1.5, 1.5]$

9. (Node2vec) Consider an undirected graph G given in the following figure. A random walk has just moved from node A to node B , with $p = 2$ and $q = 0.5$.




Figure 1. Graph G

The unnormalized transition probability from B to a next node x is $\pi_{Bx} = \alpha_{pq}(A, x)$, where $\alpha = 1/p$ if $d(A, x) = 0$, $\alpha = 1$ if $d(A, x) = 1$, and $\alpha = 1/q$ if $d(A, x) = 2$.

- Explain BFS and DFS, and explain which one a small q makes the walk behave like.
- List all possible next nodes from node B and calculate π_{Bx} .
- Calculate the normalized transition probabilities.

Solution

ENG)

- Write these three sentences.
 - BFS explores the immediate neighbourhood first, so the embedding captures the structural role of a node.
 - DFS moves far away from the source, so the embedding captures community membership.
 - A small q makes $1/q$ large, so the walk moves outward and behaves like DFS.

The previous node is A and the walk is now at B . All distances are measured from A .

A : $x = A$ is the previous node, $d = 0$, $\alpha = 1/p = 0.5$

C : x is also a neighbour of A , $d = 1$, $\alpha = 1 = 1$

D : x is two hops from A , $d = 2$, $\alpha = 1/q = 2$

Sum them: $0.5 + 1 + 2 = 3.5$

b) $\pi(B \rightarrow A) = 0.5$, $\pi(B \rightarrow C) = 1$, $\pi(B \rightarrow D) = 2$

c) $P(A) = 0.1429$, $P(C) = 0.2857$, $P(D) = 0.5714$

10. (LINE) Consider an undirected graph G with node embeddings

$$u_A = 1, u_B = 0.5, u_C = -0.5, u_D = -1$$

and context embeddings

$$u'_A = -0.5, u'_B = -1, u'_C = 0.5, u'_D = 1.$$

The second-order proximity is $p_2(j|x) = \exp(u'_j u_x) / \sum_k \exp(u'_k u_x)$. Given $e^{0.25} = 1.28$, $e^{0.5} = 1.65$, $e^{-0.25} = 0.78$, $e^{-0.5} = 0.61$.

- What is second-order proximity, and how does it differ from first-order proximity?
- Calculate $p_2(A|B)$, $p_2(B|B)$, $p_2(C|B)$ and $p_2(D|B)$.

Solution

ENG)

- Write these three sentences.

- First-order proximity is the similarity between two nodes that are directly connected by an edge.
- Second-order proximity is the similarity between two nodes that share many neighbours, even without an edge.
- So LINE gives each node two vectors: its own embedding and its context embedding.

The base node is B with $u(B) = 0.5$. The other side uses the context vector u' .

$$A: u'(A) u(B) = -0.5 \times 0.5 = -0.25, \exp = 0.78$$

$$B: u'(B) u(B) = -1 \times 0.5 = -0.5, \exp = 0.61$$

$$C: u'(C) u(B) = 0.5 \times 0.5 = 0.25, \exp = 1.28$$

$$D: u'(D) u(B) = 1 \times 0.5 = 0.5, \exp = 1.65$$

The denominator is the sum of all four: $0.78 + 0.61 + 1.28 + 1.65 = 4.32$

$$b) p_2(A|B) = 0.1806, p_2(B|B) = 0.1412, p_2(C|B) = 0.2963, p_2(D|B) = 0.3819$$

11. (GCN) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 1 & 1 \\ 1 & 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} -1 \\ -3 \\ -2 \\ 2 \\ -2 \end{bmatrix}$$

Assume that the hidden layer of a GCN model of all nodes at layer (k) can be calculated as:

$$H^{(k)} = \sigma(\tilde{A} \cdot H^{(k-1)}),$$

where $\tilde{A} = D^{-1}A$ and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- Explain how information is propagated in a GCN layer.
- Calculate $\tilde{A}$.
- Calculate the output of the GCN model at layer $k = 1$.

Solution

ENG)

- Write these three sentences.
 - In one layer each node aggregates the features of its direct neighbours and takes the average.
 - The row-normalized matrix $\tilde{A}$ equals $D^{-1}A$, which performs this averaging.
 - After k layers a node has received information from all nodes within k hops.

Degrees: $d(A)=4$, $d(B)=2$, $d(C)=2$, $d(D)=2$, $d(E)=2$

Divide each row by its degree. $\tilde{A} = [0, 0.25, 0.25, 0.25, 0.25]; [0.5, 0, 0, 0, 0.5]; [0.5, 0, 0, 0.5, 0]; [0.5, 0, 0.5, 0, 0]; [0.5, 0.5, 0, 0, 0]$ (each row sums to 1)

b) $\tilde{A} = [0, 0.25, 0.25, 0.25, 0.25]; [0.5, 0, 0, 0, 0.5]; [0.5, 0, 0, 0.5, 0]; [0.5, 0, 0.5, 0, 0]; [0.5, 0.5, 0, 0, 0]$

Multiply $\tilde{A}$ by H^0 . Each entry is the mean of that node's neighbours.

A: neighbours B, C, D, E. $(-3 + -2 + 2 + -2)/4 = -1.25$

B: neighbours A, E. $(-1 + -2)/2 = -1.5$

C: neighbours A, D. $(-1 + 2)/2 = 0.5$

D: neighbours A, C. $(-1 + -2)/2 = -1.5$

E: neighbours A, B. $(-1 + -3)/2 = -2$

ReLU clips the negatives: $-1.25 \rightarrow 0$, $-1.5 \rightarrow 0$, $-1.5 \rightarrow 0$, $-2 \rightarrow 0$

c) $H^1 = [0, 0, 0.5, 0, 0]$

12. (GraphSAGE) Consider an undirected graph G of five nodes given in the following figure. Each node has initial features that are the numbers standing next to it. According to the GraphSAGE model with an AGGREGATE that is a MEAN function, the feature of a node i at layer k can be updated as:

$$h_{N(i)}^{(k)} = \text{MEAN}(\{ h_u^{(k-1)}, \forall u \in N(i) \})$$

$$h_i^{(k)} = \text{ReLU}(h_i^{(k-1)} \parallel h_{N(i)}^{(k)})$$

where $\parallel$ is a concatenation and $\text{ReLU}(x) = \max(0, x)$.




Figure. Graph G

- What is the main idea behind GraphSAGE, and how does it differ from traditional GCNs?
- Calculate the feature of each node at $k = 1$.
- Calculate a graph-level embedding h_G by using a 'Mean' global pooling when $k = 1$.

Solution

ENG)

a) Write these three sentences.

- GraphSAGE learns an aggregator function instead of a fixed embedding for each node.
- It samples a fixed number of neighbours, so it can be trained in mini-batches on very large graphs.
- A GCN is transductive, but GraphSAGE is inductive, so it can embed a node that was never seen in training.

First compute the neighbour mean. The node itself is not included.

A: neighbours C, D, E. $(-2 + -1 + 4)/3 = 0.3333$

B: neighbours C, D. $(-2 + -1)/2 = -1.5$

C: neighbours A, B. $(-2 + -1)/2 = -1.5$

D: neighbours A, B, E. $(-2 + -1 + 4)/3 = 0.3333$

E: neighbours A, D. $(-2 + -1)/2 = -1.5$

Concatenate the node's own value with the neighbour mean. This is not an addition.

ReLU clips the negatives: $-2 \rightarrow 0$, $-1 \rightarrow 0$, $-1.5 \rightarrow 0$, $-2 \rightarrow 0$, $-1.5 \rightarrow 0$, $-1 \rightarrow 0$, $-1.5 \rightarrow 0$

The graph embedding is the mean per position: first $-2 / 5$, second $-3.8333 / 5$

b) $h(A) = [0, 0.3333]$, $h(B) = [0, 0]$, $h(C) = [0, 0]$, $h(D) = [0, 0.3333]$, $h(E) = [4, 0]$

c) $h_G = [0.8, 0.1333]$

13. (ClusterGCN) Consider an undirected graph G of six nodes given in the following figure. The graph G contains two clusters G_1 and G_2 . Each node has initial features that are the numbers standing next to it. According to the ClusterGCN model, the feature of a node i at layer k can be updated as:

$$h_{N(i)}^{(k)} = \text{MEAN}(\{ h_u^{(k-1)}, \forall u \in N(i), G_u = G_i \})$$

$$h_i^{(k)} = \text{ReLU}(h_i^{(k-1)} \parallel h_{N(i)}^{(k)})$$

where $\parallel$ is a concatenation and $\text{ReLU}(x) = \max(0, x)$.




Figure. Graph G

- What is mini-batch training, and why does ClusterGCN need it?
- Calculate the output representations of all nodes at layer $k = 1$.

Solution

ENG)

- Write these three sentences.
 - Mini-batch training updates the parameters using a small subset of the data at a time.
 - A large graph does not fit in memory, so ClusterGCN partitions the graph into densely connected clusters.
 - A good partition keeps most edges inside a batch, so little information is lost at the boundary.

Edges that cross a group are treated as absent. Cross them out first.

A: no neighbour inside the group, so the mean is 0.

B: neighbours C. $(-2)/1 = -2$

C: neighbours B. $(2)/1 = 2$

D: neighbours E. $(2)/1 = 2$

E: neighbours D. $(-1)/1 = -1$

F: no neighbour inside the group, so the mean is 0.

Concatenate the neighbour mean after the node's own value, then apply ReLU to both entries.

b) $h(A) = [0, 0]$, $h(B) = [2, 0]$, $h(C) = [0, 2]$, $h(D) = [0, 2]$, $h(E) = [2, 0]$, $h(F) = [0, 0]$

14. (GraphSAINT) Consider an undirected graph G of six nodes given in the following figure. The graph G has a sampled subgraph G_1 . Each node has initial features that are the numbers standing next to it. According to the GraphSAINT model, the feature of a node i at layer k can be updated as:

$$h_{N(i)}^{(k)} = \text{MEAN}(\{ h_u^{(k-1)}, \forall u \in N(i), G_u = G_i \})$$

$$h_i^{(k)} = \text{ReLU}(h_i^{(k-1)} \parallel h_{N(i)}^{(k)})$$

where $\parallel$ is a concatenation and $\text{ReLU}(x) = \max(0, x)$.




Figure. Graph G

- How does GraphSAINT's sampling differ from node-wise or layer-wise sampling?
- Calculate the output representations of nodes A, B, C and D at layer $k = 1$.

Solution

ENG)

- Write these three sentences.
 - GraphSAGE and FastGCN sample while building the layers, so sampling is repeated at every layer.
 - GraphSAINT samples a subgraph first and then builds the full GNN on that subgraph.
 - So the cost does not grow with depth, and normalization coefficients correct the sampling bias.

Edges that leave the sampled subgraph are discarded.

A: neighbours D. $(2)/1 = 2$

B: neighbours D. $(2)/1 = 2$

C: neighbours D. $(2)/1 = 2$

D: neighbours A, B, C. $(4 + -2 + 1)/3 = 1$

Concatenate the neighbour mean after the node's own value, then apply ReLU.

b) $h(A) = [4, 2]$, $h(B) = [0, 2]$, $h(C) = [1, 2]$, $h(D) = [2, 1]$

15. (JK Network) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 0 & 1 & 1 \\ 1 & 0 & 1 & 1 & 0 \\ 0 & 1 & 0 & 0 & 1 \\ 1 & 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 2 \\ 0 \\ 0 \\ -2 \\ 3 \end{bmatrix}$$

Assume that the output of a JK network model of all nodes at layer (k) can be calculated as:

$$H^{(k)} = \max(\sigma(\tilde{A} H^{(0)}), \sigma(\tilde{A} H^{(1)}), \dots, \sigma(\tilde{A} H^{(k-1)}))$$

where $\tilde{A} = D^{-1}A$ and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- Explain how Jumping Knowledge Networks help mitigate the over-smoothing problem.
- Calculate $\tilde{A}$ and $H^{(1)}$.
- Calculate the output representations at layer $k = 2$.

Solution

ENG)

a) Write these three sentences.

- In a deep GCN every node repeatedly averages its neighbours, so all representations become similar. This is over-smoothing.
- JKNet keeps the representation of every layer and combines them at the end by max, concatenation or an LSTM.
- The best depth is different for each node, so each node can select the depth that suits it.

Degrees: $d(A)=3$, $d(B)=3$, $d(C)=2$, $d(D)=2$, $d(E)=2$

Divide each row by its degree. $\text{Atilde} = [0, 0.3333, 0, 0.3333, 0.3333]; [0.3333, 0, 0.3333, 0.3333, 0]; [0, 0.5, 0, 0, 0.5]; [0.5, 0.5, 0, 0, 0]; [0.5, 0, 0.5, 0, 0]$ (each row sums to 1)

b) $\text{Atilde} = [0, 0.3333, 0, 0.3333, 0.3333]; [0.3333, 0, 0.3333, 0.3333, 0]; [0, 0.5, 0, 0, 0.5]; [0.5, 0.5, 0, 0, 0]; [0.5, 0, 0.5, 0, 0]$

First build $H_1 = \text{ReLU}(\text{Atilde } H_0)$

A: neighbours B, D, E. $(0 + -2 + 3)/3 = 0.3333$

B: neighbours A, C, D. $(2 + 0 + -2)/3 = 0$

C: neighbours B, E. $(0 + 3)/2 = 1.5$

D: neighbours A, B. $(2 + 0)/2 = 1$

E: neighbours A, C. $(2 + 0)/2 = 1$

No negative value, so ReLU changes nothing.

$H_1 = [0.3333, 0, 1.5, 1, 1]$

For $k = 2$ there are two terms. First: $\text{ReLU}(\text{Atilde } H_0) = [0.3333, 0, 1.5, 1, 1]$

Second: $\text{ReLU}(\text{Atilde } H_1) = [0.6667, 0.9444, 0.5, 0.1667, 0.9167]$

The max is taken element-wise, not by choosing one whole vector.

c) $H2 = [0.6667, 0.9444, 1.5, 1, 1]$

16. (GCNII) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 0 & 1 & 0 \\ 1 & 0 & 1 & 1 & 0 \\ 0 & 1 & 0 & 0 & 1 \\ 1 & 1 & 0 & 0 & 1 \\ 0 & 0 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 3 \\ 4 \\ -2 \\ 3 \\ -1 \end{bmatrix}$$

Assume that the output of a GCNII model of all nodes at layer (k) can be calculated as:

$$H^{(k)} = \sigma[((1 - \beta) I_n) \cdot ((1 - \alpha) \tilde{A} \cdot H^{(k-1)} + \alpha H^{(0)})]$$

where $\tilde{A} = D^{-1}A$, I_n is the identity matrix, $\alpha = 0.5$, $\beta = 0.5$, and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- What are the two major components introduced in GCNII to overcome the depth limitation?
- Calculate $\tilde{A}$.
- Calculate the output representations at layer $k = 1$.

Solution

ENG)

- Write these three sentences.
 - The two components are the initial residual connection and the identity mapping.
 - The initial residual connection adds a fraction alpha of the input H zero back at every layer.
 - The identity mapping keeps the weight matrix close to the identity, so a deep model is not destroyed.

Degrees: $d(A)=2$, $d(B)=3$, $d(C)=2$, $d(D)=3$, $d(E)=2$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0, 0.5, 0]; [0.3333, 0, 0.3333, 0.3333, 0]; [0, 0.5, 0, 0, 0.5]; [0.3333, 0.3333, 0, 0, 0.3333]; [0, 0, 0.5, 0.5, 0]$ (each row sums to 1)

b) $\tilde{A} = [0, 0.5, 0, 0.5, 0]; [0.3333, 0, 0.3333, 0.3333, 0]; [0, 0.5, 0, 0, 0.5]; [0.3333, 0.3333, 0, 0, 0.3333]; [0, 0, 0.5, 0.5, 0]$

Start with the inner bracket. $\tilde{A} H^0 = [3.5, 1.3333, 1.5, 2, 0.5]$

Multiply by $(1-\alpha) = 0.5$: $[1.75, 0.6667, 0.75, 1, 0.25]$

Add $\alpha H^0 = 0.5 \times [3, 4, -2, 3, -1] = [1.5, 2, -1, 1.5, -0.5]$

Inner bracket = $[3.25, 2.6667, -0.25, 2.5, -0.25]$

Multiply by the outer $(1-\beta) = 0.5$. Dropping this factor scales the answer.

ReLU clips the negatives: $-0.125 \rightarrow 0, -0.125 \rightarrow 0$

c) $H^1 = [1.625, 1.3333, 0, 1.25, 0]$

17. (DeepGCN) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 0 & 1 \\ 1 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 4 \\ -2 \\ 0 \\ -2 \end{bmatrix}$$

Assume that the output of a DeepGCN model with a residual connection at layer (k) is:

$$H^{(k)} = \sigma(\tilde{A} \cdot H^{(k-1)}) + H^{(k-1)}$$

where $\tilde{A} = D^{-1}A$ and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- Explain why a residual connection allows a GCN to be made deeper.
- Calculate the output representations at layer $k = 1$ and $k = 2$.
- Compare $H^{(1)}$ with the output of a plain GCN without the residual connection.

Solution

ENG)

- Write these three sentences.
 - Deep GCNs suffer from vanishing gradients and over-smoothing, so adding layers makes the model worse.
 - A residual connection passes the input forward unchanged, so each layer only has to learn the change.
 - If no change is useful the layer can output zero, and gradients also flow back through the skip path.

Degrees: $d(A)=2$, $d(B)=3$, $d(C)=2$, $d(D)=3$

Divide each row by its degree. Atilde = $[0, 0.5, 0, 0.5]$; $[0.3333, 0, 0.3333, 0.3333]$; $[0, 0.5, 0, 0.5]$; $[0.3333, 0.3333, 0.3333, 0]$ (each row sums to 1)

Atilde $H_0 = [-2, 0.6667, -2, 0.6667]$

Apply ReLU first: $[0, 0.6667, 0, 0.6667]$

Then add H_0 . The addition sits outside the ReLU, so negatives can survive.

For $k = 2$, feed H_1 back in and add H_1 , not H_0 .

Atilde $H_1 = [-1.3333, 0.8889, -1.3333, 0.8889]$, after ReLU $[0, 0.8889, 0, 0.8889]$

b) $H_1 = [4, -1.3333, 0, -1.3333]$

$H_2 = [4, -0.4444, 0, -0.4444]$

c) Plain GCN without residual: $H_1 = [0, 0.6667, 0, 0.6667]$ (the node's own value and negatives disappear)

5회차

1. (Relational Classification) Consider an undirected graph G of five nodes given in the following figure. Nodes 1 and 5 are labelled ($P(Y_1) = 1, P(Y_5) = 0$) and the others are unlabelled and initialized to 0.5. The relational classifier updates every unlabelled node as:

$$P(Y_v = c) = (\sum_{(v,u) \in E} A_{v,u} P(Y_u = c)) / (\sum_{(v,u) \in E} A_{v,u})$$




Figure. Graph G

- Explain the key assumption behind relational classification, and state one limitation.
- Calculate $P(Y_v)$ for every unlabelled node after the 1st iteration.
- Calculate $P(Y_v)$ after the 2nd iteration, using the previous iteration's values only.

Solution

ENG)

- Write these three sentences.
 - Relational classification assumes homophily: connected nodes tend to share the same label.
 - It repeatedly sets each unlabelled node to the weighted average of its neighbours' class probabilities.
 - Its limitation is that it uses only the graph structure and ignores node features.

Write the initial values: node 1 is 1, node 5 is 0, the rest start at 0.5.

First iteration. Only the unlabelled nodes are replaced by the mean of their neighbours.

Node 1 is labelled, so it never changes. $P = 1$

Node 2: neighbours 1, 3, 4. $(1 + 0.5 + 0.5)/3 = 0.6667$

Node 3: neighbours 2, 5. $(0.5 + 0)/2 = 0.25$

Node 4: neighbours 1, 2. $(1 + 0.5)/2 = 0.75$

Node 5 is labelled, so it never changes. $P = 0$

b) $P(Y_1) = 1, P(Y_2) = 0.6667, P(Y_3) = 0.25, P(Y_4) = 0.75, P(Y_5) = 0$

Second iteration. Use only the values from the first iteration, not the ones just computed.

Node 1 is labelled, so it never changes. $P = 1$

Node 2: neighbours 1, 3, 4. $(1 + 0.25 + 0.75)/3 = 0.6667$

Node 3: neighbours 2, 5. $(0.6667 + 0)/2 = 0.3333$

Node 4: neighbours 1, 2. $(1 + 0.6667)/2 = 0.8333$

Node 5 is labelled, so it never changes. $P = 0$

c) $P(Y_1) = 1, P(Y_2) = 0.6667, P(Y_3) = 0.3333, P(Y_4) = 0.8333, P(Y_5) = 0$

2. (FastGCN) Consider the undirected graph G with adjacency matrix A and node feature vector $h^{(0)}$ below. FastGCN replaces the full neighbourhood sum of a GCN layer by a layer-wise Monte-Carlo estimate:

$$A = \begin{bmatrix} 0 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 & 1 \\ 1 & 0 & 1 & 0 & 1 \\ 0 & 1 & 1 & 1 & 0 \end{bmatrix}$$

$$h^{(0)} = \begin{bmatrix} -2 \\ 4 \\ 4 \\ -1 \\ -2 \end{bmatrix}$$

$$h^{(l+1)}(v) = \sigma \left(\frac{1}{t_l} \sum_{j=1}^{t_l} \tilde{A}(v, u_j) h^{(l)}(u_j) \right)$$

where $\tilde{A} = D^{-1}A$ and the t_l nodes are drawn once per layer from $q(u) \propto \|\tilde{A}(:, u)\|^2$, with $\text{ReLU}(x) = \max(0, x)$.

- How does layer-wise sampling differ from the node-wise sampling of GraphSAGE?
- Calculate the importance distribution $q(u)$ over the five nodes.
- Suppose the sampled set is $\{B, E\}$ ($t_l = 2$). Calculate the estimate of $h^{(1)}(A)$.

Solution

ENG)

a) Write these three sentences.

- In node-wise sampling, each node samples its own neighbours, so the cost grows exponentially with depth.
- FastGCN samples a fixed-size set of nodes once per layer, shared by all nodes in that layer.
- Therefore the cost grows only linearly with depth instead of exponentially.

Degrees: $d(A)=2$, $d(B)=1$, $d(C)=3$, $d(D)=3$, $d(E)=3$

Divide each row by its degree. $\tilde{A} = [0, 0, 0.5, 0.5, 0]; [0, 0, 0, 0, 1]; [0.3333, 0, 0, 0.3333, 0.3333]; [0.3333, 0, 0.3333, 0, 0.3333]; [0, 0.3333, 0.3333, 0.3333, 0]$ (each row sums to 1)

$q(u)$ uses the squared norm of column u of $\tilde{A}$, not of a row.

Column A: squared sum = 0.2222

Column B: squared sum = 0.1111

Column C: squared sum = 0.4722

Column D: squared sum = 0.4722

Column E: squared sum = 1.2222

Divide by the total 2.5 to turn them into probabilities.

b) $q(A) = 0.0889$, $q(B) = 0.0444$, $q(C) = 0.1889$, $q(D) = 0.1889$, $q(E) = 0.4889$

For (c) look only at the sampled nodes: read $\tilde{A}(A, u)$ for each.

$\tilde{A}(A, B) h(B) = 0 \times 4 = 0$

$\tilde{A}(A, E) h(E) = 0 \times -2 = 0$

Add them and divide by the number sampled (2), not by the degree: $(0 + 0)/2 = 0$

No negative value, so ReLU changes nothing.

c) $h^1(A) = 0$

3. (ASGCN) Consider the undirected graph G with adjacency matrix A and node feature vector $h^{(0)}$ below. ASGCN rewrites a GCN layer as an expectation:

$$A = \begin{bmatrix} 0 & 0 & 0 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 0 & 1 & 0 \\ 1 & 1 & 1 & 0 & 1 \\ 1 & 0 & 0 & 1 & 0 \end{bmatrix}$$

$$h^{(0)} = \begin{bmatrix} 4 \\ 4 \\ 1 \\ 2 \\ 3 \end{bmatrix}$$

$$h^{(l+1)}(v_i) = \sigma(W^{(l)} N(v_i) E_{p(u_j | v_i)}[h^{(l)}(u_j)])$$

where $\tilde{A} = D^{-1}A$, $N(v_i) = \sum_j \tilde{A}_{ij}$, $p(u_j | v_i) = \tilde{A}_{ij} / N(v_i)$, $W^{(l)} = 1$ and $\text{ReLU}(x) = \max(0, x)$.

- What does the proposal distribution depend on in ASGCN that it does not in FastGCN?
- Calculate $N(v_i)$ and $p(u_j | v_i)$ for the parent node $v_i = E$.
- Calculate $E_{p(u_j | E)}[h^{(0)}(u_j)]$ and the resulting $h^{(1)}(E)$.

Solution

ENG)

- Write these three sentences.

- In FastGCN the sampling distribution is fixed in advance and does not depend on the parent nodes.
- In ASGCN the proposal distribution is conditioned on the nodes of the upper layer.
- This makes almost every sampled node useful and reduces the variance of the estimate.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=2$, $d(D)=4$, $d(E)=2$

Divide each row by its degree. $\tilde{A} = [0, 0, 0, 0.5, 0.5]; [0, 0, 0.5, 0.5, 0]; [0, 0.5, 0, 0.5, 0]; [0.25, 0.25, 0.25, 0, 0.25]; [0.5, 0, 0, 0.5, 0]$ (each row sums to 1)

The parent node is E, so only its row of $\tilde{A}$ matters.

$N(E) = \text{sum of that row} = 1$. With row normalisation it is 1.

$$p(A|E) = \tilde{A}_{EA} / N = 0.5$$

$$p(D|E) = \tilde{A}_{ED} / N = 0.5$$

$$\text{b) } N(E) = 1, p(A|E) = 0.5, p(D|E) = 0.5$$

The expectation is the sum of probability times value: $0.5 \times 4 + 0.5 \times 2 = 3$

$N = 1$ and $W = 1$, so the value is unchanged. It equals the plain neighbour mean.

No negative value, so ReLU changes nothing.

$$\text{c) } E = 3, h_1(E) = 3$$

4. (APPNP) Given a graph with an adjacency matrix A and a prediction matrix H produced by an MLP as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{bmatrix}$$

$$H = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$$

APPNP propagates the predictions with a teleport probability α :

$$Z^{(0)} = H, \quad Z^{(k+1)} = (1 - \alpha) \tilde{A} Z^{(k)} + \alpha H$$

where $\tilde{A} = D^{-1}A$ and $\alpha = 0.2$. No nonlinearity is applied between steps.

- Explain the teleport mechanism and why it prevents over-smoothing.
- Calculate $\tilde{A}$ and $Z^{(1)}$.
- Calculate $Z^{(2)}$.

Solution

ENG)

a) Write these three sentences.

- At every propagation step the walk returns to the root node with probability alpha. This is the teleport mechanism.
- Because a fixed fraction of the mass is pulled back at every step, the distribution stays local.
- So we can use many propagation steps without over-smoothing, and propagation has no extra parameters.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=2$, $d(D)=2$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0.5, 0]; [0.5, 0, 0, 0.5]; [0.5, 0, 0, 0.5]; [0, 0.5, 0.5, 0]$ (each row sums to 1)

b) $\tilde{A} = [0, 0.5, 0.5, 0]; [0.5, 0, 0, 0.5]; [0.5, 0, 0, 0.5]; [0, 0.5, 0.5, 0]$

Start from $Z_0 = H = [0, 1, 0, 0]$

Propagate one step: $\tilde{A} Z_0 = [0.5, 0, 0, 0.5]$

Multiply by $(1-\alpha) = 0.8$: $[0.4, 0, 0, 0.4]$

Add $\alpha H = [0, 0.2, 0, 0]$. This is the teleport term.

$$Z_1 = [0.4, 0.2, 0, 0.4]$$

One more step: $\tilde{A} Z_1 = [0.1, 0.4, 0.4, 0.1]$

Multiply by $(1-\alpha)$ and add αH **again**. It is added at every step.

c) $Z_2 = [0.08, 0.52, 0.32, 0.08]$

5. (GPR-GNN) Given a graph with an adjacency matrix A and hidden features $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 \\ 1 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} -1 \\ 1 \\ -2 \\ 2 \end{bmatrix}$$

GPR-GNN separates feature extraction from topological propagation:

$$H^{(k)} = \tilde{A} H^{(k-1)}, \quad Z = \sum_{k=0}^K \gamma_k H^{(k)}$$

where $\tilde{A} = D^{-1}A$, $K = 2$, $\gamma_0 = 0.7$, $\gamma_1 = 0.2$, $\gamma_2 = -0.3$.

- What can GPR-GNN do that APPNP cannot, and why do negative γ_k matter?
- Calculate $H^{(1)}$ and $H^{(2)}$.
- Calculate Z .

Solution

ENG)

- Write these three sentences.
 - APPNP fixes the coefficients to alpha times (1-alpha) to the power k, which are always positive and decreasing.
 - GPR-GNN learns the weights, and they are allowed to be negative.
 - A negative weight subtracts that hop, so the model can act as a high-pass filter on heterophilic graphs.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=3$, $d(D)=1$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0.5, 0]; [0.5, 0, 0.5, 0]; [0.3333, 0.3333, 0, 0.3333]; [0, 0, 1, 0]$
(each row sums to 1)

H_0 stays as it is: $[-1, 1, -2, 2]$

$H_1 = \tilde{A} H_0 = [-0.5, -1.5, 0.6667, -2]$

$H_2 = \tilde{A} H_1 = [-0.4167, 0.0833, -1.3333, 0.6667]$

b) $H_1 = [-0.5, -1.5, 0.6667, -2]$, $H_2 = [-0.4167, 0.0833, -1.3333, 0.6667]$

Now take the weighted sum with gamma. There is no ReLU during propagation.

$\gamma_0 H_0 = 0.7 \times [-1, 1, -2, 2] = [-0.7, 0.7, -1.4, 1.4]$

$\gamma_1 H_1 = 0.2 \times [-0.5, -1.5, 0.6667, -2] = [-0.1, -0.3, 0.1333, -0.4]$

$\gamma_2 H_2 = -0.3 \times [-0.4167, 0.0833, -1.3333, 0.6667] = [0.125, -0.025, 0.4, -0.2]$

γ_2 is negative, so that hop is subtracted. Dropping the sign gives a wrong answer.

Add the three vectors element-wise.

c) $Z = [-0.675, 0.375, -0.8667, 0.8]$

6. (MixHop) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 3 \\ -3 \\ 1 \\ 0 \end{bmatrix}$$

A MixHop layer concatenates several powers of the adjacency matrix:

$$H^{(i+1)} = \parallel_{j \in P} \sigma(\hat{A}^j H^{(i)} W_j^{(i)})$$

where $P = \{0, 1, 2\}$, $\hat{A} = D^{-1}A$, $W_j^{(i)} = 1$ and $\text{ReLU}(x) = \max(0, x)$.

- State Theorem 1 of MixHop about the two-hop Delta Operator.
- Calculate $\hat{A}^0 H^{(0)}$, $\hat{A}^1 H^{(0)}$ and $\hat{A}^2 H^{(0)}$.
- Calculate $H^{(1)}$ and the Delta Operator $\sigma(\hat{A} H^{(0)}) - \sigma(\hat{A}^2 H^{(0)})$.

Solution

ENG)

a) Write these three sentences.

- Theorem 1 says that a vanilla GCN cannot represent the two-hop Delta Operator, even with infinitely many layers.
- Stacking GCN layers collapses to $\hat{A}$ to the power 1 times W , so the intermediate hops are never kept separately.
- MixHop computes each power separately and concatenates them, so a later layer can subtract one from another.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=2$, $d(D)=2$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0.5, 0]; [0.5, 0, 0, 0.5]; [0.5, 0, 0, 0.5]; [0, 0.5, 0.5, 0]$ (each row sums to 1)

$\hat{A}^0$ is the identity, so it is just H^0 : $[3, -3, 1, 0]$

$\hat{A}^1 H^0 = [-1, 1.5, 1.5, -1]$ (one hop, the neighbour mean)

$\hat{A}^2 H^0 = \tilde{A}(\tilde{A} H^0) = [1.5, -1, -1, 1.5]$ (two hops)

ReLU is applied to each block separately, not after concatenation.

After ReLU: $[3, 0, 1, 0] \parallel [0, 1.5, 1.5, 0] \parallel [1.5, 0, 0, 1.5]$

The delta operator subtracts the two-hop block from the one-hop block.

b) $A^0 H = [3, -3, 1, 0]$, $A^1 H = [-1, 1.5, 1.5, -1]$, $A^2 H = [1.5, -1, -1, 1.5]$

c) $H^1 =$ concatenation of the three blocks: $[3, 0, 1, 0] \parallel [0, 1.5, 1.5, 0] \parallel [1.5, 0, 0, 1.5]$

Delta = $[-1.5, 1.5, 1.5, -1.5]$

7. (**PairNorm**) After one GCN layer the node features of a four-node graph are

$$X = \begin{bmatrix} -2 & 4 \\ 1 & -1 \\ -1 & 1 \\ -1 & 3 \end{bmatrix}$$

PairNorm normalizes them with a centering step and a scaling step:

$$x_i^c = x_i - (1/n) \sum_i x_i \quad (\text{Center})$$

$$x_i = s \cdot x_i^c / \text{sqrt}((1/n) \sum_i \|x_i^c\|_2^2) \quad (\text{Scale})$$

with $s = 1$ and $n = 4$.

- Explain node-wise and feature-wise over-smoothing, and what PairNorm keeps constant.
- Calculate the centered features.
- Calculate the scaling denominator and the final x_i .

Solution

ENG)

- Write these three sentences.
 - Node-wise over-smoothing means the features of connected nodes become too similar.
 - Feature-wise over-smoothing means all node features become indistinguishable across the whole graph.
 - PairNorm keeps the total pairwise squared distance between node representations constant across layers.

First take the mean over the nodes, separately for each feature column.

Column 1: $(-2 + 1 + -1 + -1)/4 = -0.75$

Column 2: $(4 + -1 + 1 + 3)/4 = 1.75$

Subtract the mean from every row. Check: each column now sums to zero.

Node 1: squared norm = 6.625

Node 2: squared norm = 10.625

Node 3: squared norm = 0.625

Node 4: squared norm = 1.625

Mean = $(6.625 + 10.625 + 0.625 + 1.625)/4 = 4.875$

Take the square root: $\text{sqrt}(4.875) = 2.2079$. Divide every row by this single number.

b) Mean = $[-0.75, 1.75]$, centered = $[-1.25, 2.25]$; $[1.75, -2.75]$; $[-0.25, -0.75]$; $[-0.25, 1.25]$

c) Scaling denominator = 2.2079, $x = [-0.5661, 1.019]$; $[0.7926, -1.2455]$; $[-0.1132, -0.3397]$; $[-0.1132, 0.5661]$

8. (DropEdge) Consider the complete graph K_4 on nodes A, B, C and D with the adjacency matrix A and initial node feature matrix $H^{(0)}$ below.

$$A = \begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} -2 \\ 2 \\ 1 \\ 0 \end{bmatrix}$$

DropEdge keeps each undirected edge independently with probability $1 - p$ at every epoch:

$$H^{(1)} = \sigma((D^{(t)})^{-1} A^{(t)} H^{(0)})$$

where $A^{(t)}$ and $D^{(t)}$ are rebuilt from the surviving edges only, $p = 1/3$, and $\text{ReLU}(x) = \max(0, x)$.

- Which two problems does DropEdge target, and how does it differ from Dropout?
- Calculate the expected number of surviving edges, and $H^{(1)}$ on the full graph.
- Suppose the edges (C, D) and (A, D) are dropped. Recalculate $H^{(1)}$.

Solution

ENG)

- Write these three sentences.
 - DropEdge targets over-fitting and over-smoothing.
 - A different random subgraph is used at every epoch, which works as data augmentation on the graph.
 - Dropout zeroes entries of the feature tensor, while DropEdge removes edges of the adjacency matrix.

K_4 has 6 edges. Each survives with probability $1 - 1/3 = 2/3$, so the expectation is 4.

First the full graph: every node has degree 3.

A: neighbours B, C, D. $(2 + 1 + 0)/3 = 1$

B: neighbours A, C, D. $(-2 + 1 + 0)/3 = -0.3333$

C: neighbours A, B, D. $(-2 + 2 + 0)/3 = 0$

D: neighbours A, B, C. $(-2 + 2 + 1)/3 = 0.3333$

Now drop the two given edges and **recount the degrees**. This is where most mistakes happen.

Degrees after dropping: $d(A)=2$, $d(B)=3$, $d(C)=2$, $d(D)=1$

A: neighbours B, C. $(2 + 1)/2 = 1.5$

B: neighbours A, C, D. $(-2 + 1 + 0)/3 = -0.3333$

C: neighbours A, B. $(-2 + 2)/2 = 0$

D: neighbours B. $(2)/1 = 2$

b) Expected number of surviving edges = 4, $H_1 = [1, 0, 0, 0.3333]$

c) $H_1 = [1.5, 0, 0, 2]$

9. (Node2vec) Consider an undirected graph G given in the following figure. A random walk has just moved from node A to node B , with $p = 2$ and $q = 0.5$.




Figure 1. Graph G

The unnormalized transition probability from B to a next node x is $\pi_{Bx} = \alpha_{pq}(A, x)$, where $\alpha = 1/p$ if $d(A, x) = 0$, $\alpha = 1$ if $d(A, x) = 1$, and $\alpha = 1/q$ if $d(A, x) = 2$.

- Explain BFS and DFS, and explain which one a small q makes the walk behave like.
- List all possible next nodes from node B and calculate π_{Bx} .
- Calculate the normalized transition probabilities.

Solution

ENG)

- Write these three sentences.
 - BFS explores the immediate neighbourhood first, so the embedding captures the structural role of a node.
 - DFS moves far away from the source, so the embedding captures community membership.
 - A small q makes $1/q$ large, so the walk moves outward and behaves like DFS.

The previous node is A and the walk is now at B . All distances are measured from A .

A : $x = A$ is the previous node, $d = 0$, $\alpha = 1/p = 0.5$

C : x is also a neighbour of A , $d = 1$, $\alpha = 1 = 1$

E : x is two hops from A , $d = 2$, $\alpha = 1/q = 2$

Sum them: $0.5 + 1 + 2 = 3.5$

b) $\pi(B \rightarrow A) = 0.5$, $\pi(B \rightarrow C) = 1$, $\pi(B \rightarrow E) = 2$

c) $P(A) = 0.1429$, $P(C) = 0.2857$, $P(E) = 0.5714$

10. (LINE) Consider an undirected graph G with node embeddings

$$u_A = 1, u_B = 0.5, u_C = -0.5, u_D = -1$$

and context embeddings

$$u'_A = 0.5, u'_B = 1, u'_C = -1, u'_D = -0.5.$$

The second-order proximity is $p_2(j|x) = \exp(u'_j u_x) / \sum_k \exp(u'_k u_x)$. Given $e^{0.25} = 1.28$, $e^{0.5} = 1.65$, $e^{-0.25} = 0.78$, $e^{-0.5} = 0.61$.

- What is second-order proximity, and how does it differ from first-order proximity?
- Calculate $p_2(A|B)$, $p_2(B|B)$, $p_2(C|B)$ and $p_2(D|B)$.

Solution

ENG)

- Write these three sentences.

- First-order proximity is the similarity between two nodes that are directly connected by an edge.
- Second-order proximity is the similarity between two nodes that share many neighbours, even without an edge.
- So LINE gives each node two vectors: its own embedding and its context embedding.

The base node is B with $u(B) = 0.5$. The other side uses the context vector u' .

$$A: u'(A) u(B) = 0.5 \times 0.5 = 0.25, \exp = 1.28$$

$$B: u'(B) u(B) = 1 \times 0.5 = 0.5, \exp = 1.65$$

$$C: u'(C) u(B) = -1 \times 0.5 = -0.5, \exp = 0.61$$

$$D: u'(D) u(B) = -0.5 \times 0.5 = -0.25, \exp = 0.78$$

The denominator is the sum of all four: $1.28 + 1.65 + 0.61 + 0.78 = 4.32$

$$b) p_2(A|B) = 0.2963, p_2(B|B) = 0.3819, p_2(C|B) = 0.1412, p_2(D|B) = 0.1806$$

11. (GCN) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 0 & 0 \\ 1 & 0 & 0 & 1 & 1 \\ 1 & 0 & 0 & 1 & 1 \\ 0 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 0 \\ 4 \\ 0 \\ -1 \\ -3 \end{bmatrix}$$

Assume that the hidden layer of a GCN model of all nodes at layer (k) can be calculated as:

$$H^{(k)} = \sigma(\tilde{A} \cdot H^{(k-1)}),$$

where $\tilde{A} = D^{-1}A$ and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- Explain how information is propagated in a GCN layer.
- Calculate $\tilde{A}$.
- Calculate the output of the GCN model at layer $k = 1$.

Solution

ENG)

a) Write these three sentences.

- In one layer each node aggregates the features of its direct neighbours and takes the average.
- The row-normalized matrix $\tilde{A}$ equals $D^{-1}A$, which performs this averaging.
- After k layers a node has received information from all nodes within k hops.

Degrees: $d(A)=2$, $d(B)=3$, $d(C)=3$, $d(D)=2$, $d(E)=2$

Divide each row by its degree. $\tilde{A} = [0, 0.5, 0.5, 0, 0]; [0.3333, 0, 0, 0.3333, 0.3333]; [0.3333, 0, 0, 0.3333, 0.3333]; [0, 0.5, 0.5, 0, 0]; [0, 0.5, 0.5, 0, 0]$ (each row sums to 1)

b) $\tilde{A} = [0, 0.5, 0.5, 0, 0]; [0.3333, 0, 0, 0.3333, 0.3333]; [0.3333, 0, 0, 0.3333, 0.3333]; [0, 0.5, 0.5, 0, 0]; [0, 0.5, 0.5, 0, 0]$

Multiply $\tilde{A}$ by H^0 . Each entry is the mean of that node's neighbours.

A: neighbours B, C. $(4 + 0)/2 = 2$

B: neighbours A, D, E. $(0 + -1 + -3)/3 = -1.3333$

C: neighbours A, D, E. $(0 + -1 + -3)/3 = -1.3333$

D: neighbours B, C. $(4 + 0)/2 = 2$

E: neighbours B, C. $(4 + 0)/2 = 2$

ReLU clips the negatives: $-1.3333 \rightarrow 0$, $-1.3333 \rightarrow 0$

c) $H_1 = [2, 0, 0, 2, 2]$

12. (GraphSAGE) Consider an undirected graph G of five nodes given in the following figure. Each node has initial features that are the numbers standing next to it. According to the GraphSAGE model with an AGGREGATE that is a MEAN function, the feature of a node i at layer k can be updated as:

$$h_{N(i)}^{(k)} = \text{MEAN}(\{ h_u^{(k-1)}, \forall u \in N(i) \})$$

$$h_i^{(k)} = \text{ReLU}(h_i^{(k-1)} \parallel h_{N(i)}^{(k)})$$

where $\parallel$ is a concatenation and $\text{ReLU}(x) = \max(0, x)$.




Figure. Graph G

- What is the main idea behind GraphSAGE, and how does it differ from traditional GCNs?
- Calculate the feature of each node at $k = 1$.
- Calculate a graph-level embedding h_G by using a 'Mean' global pooling when $k = 1$.

Solution

ENG)

a) Write these three sentences.

- GraphSAGE learns an aggregator function instead of a fixed embedding for each node.
- It samples a fixed number of neighbours, so it can be trained in mini-batches on very large graphs.
- A GCN is transductive, but GraphSAGE is inductive, so it can embed a node that was never seen in training.

First compute the neighbour mean. The node itself is not included.

A: neighbours C, D, E. $(-1 + 0 + 2)/3 = 0.3333$

B: neighbours D, E. $(0 + 2)/2 = 1$

C: neighbours A, E. $(3 + 2)/2 = 2.5$

D: neighbours A, B. $(3 + 4)/2 = 3.5$

E: neighbours A, B, C. $(3 + 4 + -1)/3 = 2$

Concatenate the node's own value with the neighbour mean. This is not an addition.

ReLU clips the negatives: $-1 \rightarrow 0$

The graph embedding is the mean per position: first $8 / 5$, second $9.3333 / 5$

b) $h(A) = [3, 0.3333]$, $h(B) = [4, 1]$, $h(C) = [0, 2.5]$, $h(D) = [0, 3.5]$, $h(E) = [2, 2]$

c) $h_G = [1.8, 1.8667]$

13. (ClusterGCN) Consider an undirected graph G of six nodes given in the following figure. The graph G contains two clusters G_1 and G_2 . Each node has initial features that are the numbers standing next to it. According to the ClusterGCN model, the feature of a node i at layer k can be updated as:

$$h_{N(i)}^{(k)} = \text{MEAN}(\{ h_u^{(k-1)}, \forall u \in N(i), G_u = G_i \})$$

$$h_i^{(k)} = \text{ReLU}(h_i^{(k-1)} \parallel h_{N(i)}^{(k)})$$

where $\parallel$ is a concatenation and $\text{ReLU}(x) = \max(0, x)$.




Figure. Graph G

- What is mini-batch training, and why does ClusterGCN need it?
- Calculate the output representations of all nodes at layer $k = 1$.

Solution

ENG)

- Write these three sentences.
 - Mini-batch training updates the parameters using a small subset of the data at a time.
 - A large graph does not fit in memory, so ClusterGCN partitions the graph into densely connected clusters.
 - A good partition keeps most edges inside a batch, so little information is lost at the boundary.

Edges that cross a group are treated as absent. Cross them out first.

A: neighbours B. $(4)/1 = 4$

B: neighbours A. $(0)/1 = 0$

C: no neighbour inside the group, so the mean is 0.

D: neighbours F. $(4)/1 = 4$

E: neighbours F. $(4)/1 = 4$

F: neighbours D, E. $(2 + 0)/2 = 1$

Concatenate the neighbour mean after the node's own value, then apply ReLU to both entries.

b) $h(A) = [0, 4]$, $h(B) = [4, 0]$, $h(C) = [4, 0]$, $h(D) = [2, 4]$, $h(E) = [0, 4]$, $h(F) = [4, 1]$

14. (GraphSAINT) Consider an undirected graph G of six nodes given in the following figure. The graph G has a sampled subgraph G_1 . Each node has initial features that are the numbers standing next to it. According to the GraphSAINT model, the feature of a node i at layer k can be updated as:

$$h_{N(i)}^{(k)} = \text{MEAN}(\{ h_u^{(k-1)}, \forall u \in N(i), G_u = G_i \})$$

$$h_i^{(k)} = \text{ReLU}(h_i^{(k-1)} \parallel h_{N(i)}^{(k)})$$

where $\parallel$ is a concatenation and $\text{ReLU}(x) = \max(0, x)$.




Figure. Graph G

- How does GraphSAINT's sampling differ from node-wise or layer-wise sampling?
- Calculate the output representations of nodes A, B, C and D at layer $k = 1$.

Solution

ENG)

- Write these three sentences.
 - GraphSAGE and FastGCN sample while building the layers, so sampling is repeated at every layer.
 - GraphSAINT samples a subgraph first and then builds the full GNN on that subgraph.
 - So the cost does not grow with depth, and normalization coefficients correct the sampling bias.

Edges that leave the sampled subgraph are discarded.

A: neighbours B. $(1)/1 = 1$

B: neighbours A, C, D. $(3 + 3 + 3)/3 = 3$

C: neighbours B. $(1)/1 = 1$

D: neighbours B. $(1)/1 = 1$

Concatenate the neighbour mean after the node's own value, then apply ReLU.

b) $h(A) = [3, 1]$, $h(B) = [1, 3]$, $h(C) = [3, 1]$, $h(D) = [3, 1]$

15. (JK Network) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 0 & 1 \\ 1 & 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 & 1 \\ 0 & 0 & 1 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 1 \\ -1 \\ 4 \\ 2 \\ -1 \end{bmatrix}$$

Assume that the output of a JK network model of all nodes at layer (k) can be calculated as:

$$H^{(k)} = \max(\sigma(\tilde{A} H^{(0)}), \sigma(\tilde{A} H^{(1)}), \dots, \sigma(\tilde{A} H^{(k-1)}))$$

where $\tilde{A} = D^{-1}A$ and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- Explain how Jumping Knowledge Networks help mitigate the over-smoothing problem.
- Calculate $\tilde{A}$ and $H^{(1)}$.
- Calculate the output representations at layer $k = 2$.

Solution

ENG)

a) Write these three sentences.

- In a deep GCN every node repeatedly averages its neighbours, so all representations become similar. This is over-smoothing.
- JKNet keeps the representation of every layer and combines them at the end by max, concatenation or an LSTM.
- The best depth is different for each node, so each node can select the depth that suits it.

Degrees: $d(A)=3$, $d(B)=2$, $d(C)=3$, $d(D)=1$, $d(E)=3$

Divide each row by its degree. $\text{Atilde} = [0, 0.3333, 0.3333, 0, 0.3333]; [0.5, 0, 0, 0, 0.5]; [0.3333, 0, 0, 0.3333, 0.3333]; [0, 0, 1, 0, 0]; [0.3333, 0.3333, 0.3333, 0, 0]$ (each row sums to 1)

b) $\text{Atilde} = [0, 0.3333, 0.3333, 0, 0.3333]; [0.5, 0, 0, 0, 0.5]; [0.3333, 0, 0, 0.3333, 0.3333]; [0, 0, 1, 0, 0]; [0.3333, 0.3333, 0.3333, 0, 0]$

First build $H_1 = \text{ReLU}(\text{Atilde } H_0)$

A: neighbours B, C, E. $(-1 + 4 + -1)/3 = 0.6667$

B: neighbours A, E. $(1 + -1)/2 = 0$

C: neighbours A, D, E. $(1 + 2 + -1)/3 = 0.6667$

D: neighbours C. $(4)/1 = 4$

E: neighbours A, B, C. $(1 + -1 + 4)/3 = 1.3333$

No negative value, so ReLU changes nothing.

$H_1 = [0.6667, 0, 0.6667, 4, 1.3333]$

For $k = 2$ there are two terms. First: $\text{ReLU}(\text{Atilde } H_0) = [0.6667, 0, 0.6667, 4, 1.3333]$

Second: $\text{ReLU}(\text{Atilde } H_1) = [0.6667, 1, 2, 0.6667, 0.4444]$

The max is taken element-wise, not by choosing one whole vector.

c) $H_2 = [0.6667, 1, 2, 4, 1.3333]$

16. (GCNII) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 0 & 1 \\ 1 & 1 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 1 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} 1 \\ -3 \\ 3 \\ 4 \\ -1 \end{bmatrix}$$

Assume that the output of a GCNII model of all nodes at layer (k) can be calculated as:

$$H^{(k)} = \sigma[((1 - \beta) I_n) \cdot ((1 - \alpha) \tilde{A} \cdot H^{(k-1)} + \alpha H^{(0)})]$$

where $\tilde{A} = D^{-1}A$, I_n is the identity matrix, $\alpha = 0.5$, $\beta = 0.4$, and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- What are the two major components introduced in GCNII to overcome the depth limitation?
- Calculate $\tilde{A}$.
- Calculate the output representations at layer $k = 1$.

Solution

ENG)

- Write these three sentences.
 - The two components are the initial residual connection and the identity mapping.
 - The initial residual connection adds a fraction alpha of the input H zero back at every layer.
 - The identity mapping keeps the weight matrix close to the identity, so a deep model is not destroyed.

Degrees: $d(A)=2$, $d(B)=2$, $d(C)=3$, $d(D)=3$, $d(E)=2$

Divide each row by its degree. $\tilde{A} = [0, 0, 0.5, 0.5, 0]; [0, 0, 0.5, 0, 0.5]; [0.3333, 0.3333, 0, 0.3333, 0]; [0.3333, 0, 0.3333, 0, 0.3333]; [0, 0.5, 0, 0.5, 0]$ (each row sums to 1)

b) $\tilde{A} = [0, 0, 0.5, 0.5, 0]; [0, 0, 0.5, 0, 0.5]; [0.3333, 0.3333, 0, 0.3333, 0]; [0.3333, 0, 0.3333, 0, 0.3333]; [0, 0.5, 0, 0.5, 0]$

Start with the inner bracket. $\tilde{A} H^0 = [3.5, 1, 0.6667, 1, 0.5]$

Multiply by $(1-\alpha) = 0.5$: $[1.75, 0.5, 0.3333, 0.5, 0.25]$

Add $\alpha H^0 = 0.5 \times [1, -3, 3, 4, -1] = [0.5, -1.5, 1.5, 2, -0.5]$

Inner bracket = $[2.25, -1, 1.8333, 2.5, -0.25]$

Multiply by the outer $(1-\beta) = 0.6$. Dropping this factor scales the answer.

ReLU clips the negatives: $-0.6 \rightarrow 0$, $-0.15 \rightarrow 0$

c) $H^1 = [1.35, 0, 1.1, 1.5, 0]$

17. (DeepGCN) Given a graph with an adjacency matrix A and initial node feature matrix $H^{(0)}$ as follows:

$$A = \begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 0 \\ 1 & 1 & 0 & 0 \end{bmatrix}$$

$$H^{(0)} = \begin{bmatrix} -3 \\ 1 \\ 2 \\ 0 \end{bmatrix}$$

Assume that the output of a DeepGCN model with a residual connection at layer (k) is:

$$H^{(k)} = \sigma(\tilde{A} \cdot H^{(k-1)}) + H^{(k-1)}$$

where $\tilde{A} = D^{-1}A$ and σ is a ReLU function $\text{ReLU}(x) = \max(0, x)$.

- Explain why a residual connection allows a GCN to be made deeper.
- Calculate the output representations at layer $k = 1$ and $k = 2$.
- Compare $H^{(1)}$ with the output of a plain GCN without the residual connection.

Solution

ENG)

- Write these three sentences.
 - Deep GCNs suffer from vanishing gradients and over-smoothing, so adding layers makes the model worse.
 - A residual connection passes the input forward unchanged, so each layer only has to learn the change.
 - If no change is useful the layer can output zero, and gradients also flow back through the skip path.

Degrees: $d(A)=3$, $d(B)=3$, $d(C)=2$, $d(D)=2$

Divide each row by its degree. Atilde = $[0, 0.3333, 0.3333, 0.3333]$; $[0.3333, 0, 0.3333, 0.3333]$; $[0.5, 0.5, 0, 0]$; $[0.5, 0.5, 0, 0]$ (each row sums to 1)

Atilde $H_0 = [1, -0.3333, -1, -1]$

Apply ReLU first: $[1, 0, 0, 0]$

Then add H_0 . The addition sits outside the ReLU, so negatives can survive.

For $k = 2$, feed H_1 back in and add H_1 , not H_0 .

Atilde $H_1 = [1, 0, -0.5, -0.5]$, after ReLU $[1, 0, 0, 0]$

b) $H_1 = [-2, 1, 2, 0]$

$H_2 = [-1, 1, 2, 0]$

c) Plain GCN without residual: $H_1 = [1, 0, 0, 0]$ (the node's own value and negatives disappear)